

Food delivery & quick commerce

MUMBAI, A TUESDAY IN JULY, 8:40 PM

It has rained for three hours. In Andheri West, orders on the food app are running at twice a normal Tuesday: nobody wants to step out, and everybody wants hot food. At the same moment, half the delivery riders who were online at 7 pm have logged off. Roads are flooded, two-wheelers stall in knee-deep water, and a rider who has already earned ₹900 today would rather go home dry. On the operations dashboard, the average promised delivery time has jumped from 31 to 52 minutes, and customer cancellations are climbing every minute.

Try it: you run the delivery app for Mumbai tonight. What do you change in the next 30 minutes, and what must you not break? Think for five minutes before reading on.

How a ₹450 dinner makes (a little) money

9 pm, Bengaluru. Kavya orders a ₹450 chicken biryani on Zomato and pays ₹490 with fees. The restaurant cooks it, a rider she will never see again carries it 4 km, and it arrives in 31 minutes. Three businesses were paid out of that ₹490, and the app kept less than ₹40 of it. A decade of losses and price wars in this industry came down to that last ₹40.

₹10,769 Cr Zomato food order value, Apr–Jun 2026, up 20% on a year earlier	₹9,490 Cr Swiggy food order value in the same quarter, up 17%	₹17,132 Cr Blinkit order value in the same quarter, up 86%: now bigger than Zomato food	~6,000 dark stores across India (Apr 2026), almost none before 2021
--	---	---	---

Three customers, one app

A food app is a **three-sided marketplace**: a business that earns by matching three groups who need each other. The **customer** wants choice, speed and a fair price. The **restaurant** wants orders without paying for its own riders. The **rider** (the platform calls them delivery partners) wants steady earnings per hour. If any one side is unhappy, the other two suffer: too few riders and food goes cold; too few restaurants and customers leave; too few orders and riders log off. Follow Kavya's money in Figure 2.1.





Figure 2.1 • Who pays whom on one food order. The customer's ₹490 and the restaurant's ₹15 of ads come in; ₹445 goes straight out to the restaurant and the rider. Everything the platform spends on discounts, refunds, customer support and its own staff must come out of the ₹60 left in the middle. Illustrative order; commission of 18–25% and fees of ₹30–50 are typical in large cities (2025–26). GST ignored.


Two business models sit under the same app icon. A **marketplace (third-party, 3P) model** matches buyers with sellers who own the goods: Zomato never owns the biryani, it earns a commission. An **inventory (first-party, 1P) model** buys stock and sells it itself, keeping the whole margin and the whole risk. Indian rules forbid foreign-owned e-commerce companies from selling their own inventory, so quick-commerce apps started as marketplaces with “partner sellers” in their dark stores. In 2025 Eternal (Zomato's parent) became Indian-owned and controlled, and Blinkit began moving to owning its stock: higher margin per basket, but unsold milk and bread are now its own loss.

Marketplace (3P): the app earns a fee



Earns: commission + fees, ~15–18%
 Milk expires: the seller's loss
 Price and range: the seller decides

Inventory (1P): the app is the shop



Earns: the full product margin, ~20%+
 Milk expires: Blinkit's loss
 Price and range: Blinkit decides
 Allowed only for Indian-owned and controlled companies

Figure 2.2 · The same packet of milk, two business models. Owning the stock adds margin and control, and moves the risk of waste onto the platform. That is why fill rate and waste become a quick-commerce PM's daily numbers. *India's FDI policy for e-commerce (Press Note 2, 2018); Eternal shareholder letters, 2025–26. Margins illustrative.*

Where the minutes go

“30 minutes” and “10 minutes” sound like the same promise at different speeds. They are different machines. Break both into steps and the biggest block is in a different place (Figure 2.3).

Food order, 4 km (Kavya's biryani)




Figure 2.3 · Thirty minutes vs ten, step by step. The rider is dispatched while the food cooks so that pickup and cooking finish together; a rider who arrives early waits unpaid for time and blocks other orders. In quick commerce nobody cooks, so the only way to 10 minutes is a store about 2 km from every home. *Illustrative timings consistent with company descriptions of dispatch and dark-store picking (2024–26).*

Unit economics: one order, in rupees

Unit economics means the profit or loss on one unit, here one order, before the costs that don't grow with orders (head office, brand ads). The number that matters is **contribution margin**: what the order brings in minus the costs it directly causes. If it is negative, growth makes losses bigger. Figure 2.4 walks Kavya's order through.




Figure 2.4 · Where the ₹145 the platform earns on a ₹450 order goes. Rider pay eats more than half. Every change a PM ships (a new fee, a bigger discount, a longer ride) moves one of these bars, and the ₹37 at the bottom is what's left to fight over. *Illustrative; calibrated to Eternal's reported food-delivery adjusted EBITDA of 5.6% of order value, Apr–Jun 2026.*

Quick commerce is a different shape (Figure 2.5). The basket is bigger and, since Blinkit owns the stock, it earns a product margin rather than a commission. But it also pays for a store: rent, pickers, electricity, and the goods that expire on the shelf.





Figure 2.5 • A ₹650 quick-commerce basket. More money comes in than on a food order, and more goes out, because the platform now runs a shop. The store costs are fixed per day, so the same store at half the orders would lose about ₹30 on this basket. *Illustrative; calibrated to Swiggy Instamart reaching contribution breakeven in May 2026 at an average basket of ~₹700.*

Why density beats speed

A dark store is a small warehouse closed to walk-in shoppers, stocked with 3,000 to 10,000 items and placed so that a rider can reach most homes around it in five minutes. Its rent and staff cost the same whether it handles 300 orders a day or 1,500. So the question every quick-commerce PM lives with is how many orders sit inside that 2 km circle (Figure 2.6).

Same 2 km circle, same store cost (illustrative)



To reach 1,300 orders here the circle must grow to ~4.5 km, and the ride alone then takes 12 minutes.

Figure 2.6 • Orders per store decide the economics. A mature Blinkit store handles about ₹8.3 lakh of orders a day, roughly 1,250 baskets. Dense wards reach that; small towns can't without breaking the 10-minute promise, which is why expansion beyond the top cities is slow and costly. *Blinkit store productivity from Eternal's Apr–Jun 2026 shareholder letter; the two wards are illustrative.*

The players, mid-2026

Company	Businesses	Scale (Apr–Jun 2026 unless dated)	Where it stands
Eternal (listed 2021 as Zomato)	Zomato food, Blinkit, District (going out), Hyperpure (supplies to restaurants)	Food ₹10,769 Cr order value, 5.6% ad-adjusted EBITDA; Blinkit ₹17,132 Cr, first quarterly profit (₹102 Cr)	Leader in both; profitable group
Swiggy (listed Nov 2024)	Food, Instamart, Dineout, Scenes (events)	Food ₹9,490 Cr; Instamart ₹7,907 Cr, contribution breakeven in May 2026; net loss ₹791 Cr	Close number two in food; number three in quick commerce
Zepto (private)	Quick commerce, Zepto Café	FY26 loss about ₹5,900 Cr (press reports)	Grew fastest; IPO put on hold, Jul 2026
BigBasket (Tata)	BB Now (10–20 min), scheduled grocery	FY26 revenue ~₹10,500 Cr, loss ~₹3,200 Cr	Strong in fresh produce; late to speed
Flipkart Minutes, Amazon Now	Quick commerce inside the big e-commerce apps	Amazon Now expanding to 100 cities (2026)	Deep pockets, patient capital
magicpin on ONDC	Food delivery through the government-backed open network	~1.5 lakh orders a day (Oct 2024)	Challenger on lower commissions



Order value in one quarter, Apr–Jun 2026. Zomato and Blinkit report net order value; Swiggy reports gross order value, which is slightly higher for the same orders.

Figure 2.7 · Quick commerce overtook food in two years. Food grows about 17–20% a year; the leader in 10-minute delivery grew 86%. Profit is the other way round: food makes money, quick commerce has only just stopped losing it. *Eternal and Swiggy results for Q1 FY27 (Jul 2026).*

Turning points




Figure 2.8 · Eighteen years in eight moments. The industry moved from listing restaurants, to delivering food, to running shops. Each step raised the fixed costs and the prize. *Company filings and press coverage (Business Standard, Reuters, Inc42), 2008–2026.*

The metrics that matter

Metric	What it means, with an example	Why a PM watches it
GOV / NOV (gross or net order value)	Everything customers paid in a period; net removes discounts and some fees. Kavya's order counts ₹490 in GOV	Size and growth; the denominator for every margin
AOV (average order value)	GOV ÷ orders. Food ~₹430; quick commerce ₹650–700	Fixed rider cost spread over a bigger basket
MTU (monthly transacting users) and frequency	People who ordered at least once in the month, and orders per such user (food: 3–4)	Growth comes from more users or more orders each
Contribution margin per order	₹ left after the order's own costs: the ₹37 in Figure 2.4	Tells you whether growth makes or burns money
Adjusted EBITDA % of GOV	Profit before interest, tax, depreciation and share-based pay, as a share of order value. Zomato food: 5.6%	The number investors quote
On-time % and cancellations	Orders delivered within the promised time; orders cancelled by customer, restaurant or platform	The promise is the product; every broken one costs a refund or a customer
Fill rate	Share of items ordered that were in stock and delivered (quick commerce)	A missing item ruins a 10-minute promise
Rider orders per hour	Deliveries one rider completes in an hour online	Higher means cheaper orders and better rider earnings

Five mechanics that aren't obvious

1. Batching. One rider carries two orders from the same area. Cost per order falls by a third, but the second customer waits a few minutes longer. The algorithm decides the trade-off hundreds of times a minute, and the PM sets the rule for how much delay a batch may add.

2. Platform fees are pure margin. Swiggy started a ₹2 per-order fee in 2023; by 2025 both apps charged around ₹10–15 in big cities. At Zomato's ~8 crore food orders a month, each extra rupee is about ₹8 Cr a month, with no extra cost. The only test is whether conversion drops.

3. Ads are the second business. Restaurants and brands pay to appear higher in search and on the home screen: roughly 3–4% of order value, at almost 100% margin. That is why a sponsored listing can be more profitable than the order itself.

4. Riders are now covered by labour law. The Code on Social Security came into force on 21 November 2025 and asks platforms to pay 1–2% of annual turnover (capped at 5% of what they pay workers) into a welfare fund. Karnataka's 2025 law adds a welfare fee of 1–5% of each payout, and Rajasthan passed a similar law in 2023. Rider cost per order rises a little, permanently.

5. Restaurants push back. A Competition Commission of India (CCI) investigation report in 2024 said both apps favoured restaurants that signed exclusivity or price-parity terms (the restaurant must not sell cheaper elsewhere). Restaurants also run their own delivery and join ONDC to pay lower commissions. A commission rise that looks free on a spreadsheet can push the best restaurants away.

REAL-LIFE EXAMPLE · WHAT THE APPS DO ON A RAIN NIGHT

Back to the Mumbai monsoon. The playbook, which both big apps run in some form: a **rain surge fee** of ₹20–30 on each order, shown before checkout, which cools demand a little and pays for the next point; a **rain incentive for riders**, extra rupees per order, so fewer log off; **honest, longer promised times**, because a 52-minute promise kept beats a 30-minute promise broken; and **switching off far-away restaurants** so that riders are not sent on 8 km trips through flooded roads. What must not break: rider safety (no incentive that pays for riding through water), and trust (fees shown up front, never added after payment).

CHECK · “TEN-MINUTE DELIVERY MAKES RIDERS RACE THROUGH TRAFFIC” **OVERSTATED**

Most of the ten minutes is bought by where the store is, not by how fast the rider rides: 2 km at 25 km/h is under five minutes. The riskier part is the pressure a visible countdown creates. In January 2026, after the labour ministry raised rider safety with the platforms, Blinkit dropped “10 minutes” from its branding (press reports). The lesson for a PM: promise the customer a time, but don’t turn it into a rider’s penalty. Measure rider safety (accidents per lakh orders) as a guardrail next to delivery time.

WORKED EXAMPLE · SHOULD THE PLATFORM FEE GO FROM ₹10 TO ₹15?

Illustrative, for one city with 10 lakh orders a month. **Gain:** ₹5 × 10 lakh = ₹50 lakh a month, all margin. **Loss:** if 2% of orders don’t happen, that is 20,000 orders × ₹37 of contribution each = ₹7.4 lakh. Net: about +₹43 lakh. The fee rise pays unless more than about 13% of orders disappear (₹50 lakh ÷ ₹37 = 1.35 lakh orders). The numbers look easy, and that is the trap: the loss rarely shows up in the first week. It shows up months later as customers who order a little less often, or try the rival app whose fee is ₹5 lower. So test it in two or three cities against similar ones that keep the old fee, watch orders per customer for eight weeks, and keep the fee as a flat rupee amount shown before checkout.

KEY IDEA · THREE QUESTIONS FOR ANY ORDER

1. Who pays for this order, and who is paid?
2. Which step eats the minutes?
3. Does the next order in this store, kitchen or rider shift make the first one cheaper?

IN THE INTERVIEW · QUESTIONS THEY LOVE HERE

1. Delivery times rose in one area: diagnose it.
2. Cancellations went up: why, and what would you do?
3. Should the platform fee go from ₹10 to ₹15?
4. How many dark stores does a city need?
5. Improve Swiggy or Zomato for one kind of user.

CASE 2.1 · A VENDOR PLATFORM FOR QUICK COMMERCE**Product design**

Medium

IIM A *Product Mindset 2026–27*, pp. 269–271

Try it first. A dark store runs out of Amul butter at 8 pm. Trace back the chain of messages between the app's buyers and Amul that should have prevented it. Where would a vendor platform help most?

THE QUESTION

Our quick commerce platform works with thousands of vendors. Today, purchase orders, shipment updates and invoices flow through a mix of EDI integrations, emails and spreadsheets, making procurement inefficient. As the Product Manager, design a Vendor Management Platform that streamlines supplier collaboration and standardizes these workflows. Walk me through your approach.

A Before I begin, I'd like to understand the problem a little better. Can I clarify a few things?

- Will this platform be used only by our internal procurement teams, or will vendors also have access?
- Are we replacing existing procurement systems or building a layer that integrates with existing ERPs and EDI systems?
- What is the primary business goal - reducing operational effort, improving inventory availability, or improving vendor collaboration?

Q Good questions. The platform will be used by both procurement teams and vendors. We're not replacing ERPs like SAP. Instead, we're building a collaboration layer that integrates with existing ERP and EDI systems. The primary goal is to reduce manual operations, improve vendor communication, and ensure products reach our dark stores on time.

A Understood. I'll first identify the users and their pain points, then map the procurement workflow, identify opportunities, prioritize an MVP, and finally discuss success metrics. Does that approach sound good?

Q Yes, please go ahead.

A I see four primary stakeholders.

Procurement Managers

- Create purchase orders
- Track deliveries
- Compare vendor performance
- Resolve supply issues

Vendors

- Receive purchase orders
- Confirm quantities
- Send shipment updates
- Submit invoices
- Track payments

Warehouse Teams

- Receive shipments
- Verify quantities
- Report shortages and damages

Finance Teams

- Match invoices against purchase orders
- Process payments
- Handle disputes

I'd prioritize Procurement Managers and Vendors because they directly influence procurement efficiency and inventory availability.

Q Makes sense. What problems are they facing today?

A From the problem statement, I see a few major challenges.

For procurement teams:

- Vendor communication happens across emails, phone calls and multiple EDI channels.
- Shipment visibility is poor.
- Exceptions like delayed deliveries require constant follow-ups.
- Vendor performance is difficult to track.

For vendors:

- Different retailers require different EDI formats.

- It's difficult to know whether a purchase order was received successfully.
- Payment and invoice status isn't transparent.
- Disputes take a long time to resolve. Overall, everyone spends significant effort coordinating instead of making procurement decisions.

Q That's fair. What would you build?

A Before jumping into features, I'd prioritize them using the RICE framework, balancing business impact with implementation effort.

Using a simple RICE lens:

- **Purchase Order Hub:** Highest reach since every procurement transaction starts here, with high business impact and relatively low implementation complexity.
- **EDI Integration:** High impact but lower reach initially, as not all vendors use EDI.
- **Shipment Visibility:** Depends on purchase orders already flowing through the platform.
- **Vendor Scorecards:** Valuable, but requires historical transaction data, making it better suited after adoption.

So my launch order would remain:

1. Purchase Order Hub
2. EDI Integration & Monitoring
3. Shipment Visibility
4. Vendor Performance Dashboard

For the MVP, I'd focus on reducing manual coordination and improving visibility.

The platform would include four core capabilities.

Feature	What it is
Unified Purchase Order Hub	Procurement teams can create purchase orders through the ERP, while vendors receive them through a single portal regardless of whether they connect via EDI or manually. The platform tracks every purchase order through its lifecycle: from creation and acknowledgement to shipment, invoice status and payment status.
EDI Integration & Monitoring	The system integrates with common EDI standards and automatically validates incoming and outgoing documents. If an EDI transaction fails or contains incorrect data, both sides receive alerts instead of discovering the issue later. This reduces manual troubleshooting significantly.
Shipment Visibility Dashboard	Vendors can acknowledge purchase orders, submit Advance Ship Notices (ASN), update shipment status and expected delivery dates. Procurement teams receive real-time visibility into shipment progress.
Vendor Performance Dashboard	Every vendor receives a scorecard showing: • On-time delivery rate • Fill rate • PO acknowledgement time • Invoice accuracy • EDI transaction success rate This helps procurement teams make better sourcing decisions.

Q These all sound useful. Which one would you launch first?

A I'd prioritize:

1. Purchase Order Hub
2. EDI Integration & Monitoring
3. Shipment Visibility
4. Vendor Performance Dashboard

Purchase orders form the foundation of the procurement process. Once they're flowing reliably through the platform, shipment tracking and analytics become much easier to build.

Q What would you intentionally leave out of the MVP?

A I'd postpone:

- AI demand forecasting
- Automated vendor recommendations
- Dynamic procurement negotiations
- Contract lifecycle management
- Predictive stock replenishment

These features become much more valuable once we've collected reliable procurement and vendor performance data.

Q Good. How would you measure success?

A I'd divide metrics into four buckets.

Metric Bucket	Metric
Adoption	• % of vendors onboarded to the platform • % of purchase orders processed through the platform • % of vendors using EDI integrations • Weekly active procurement managers • Weekly active vendor users
Operational Efficiency	• Average purchase order processing time • Average PO acknowledgement time • Average time to resolve EDI failures • Number of manual follow-ups per purchase order • % of procurement workflows completed without manual intervention
Supply Reliability	• On-time delivery rate • Fill rate • Shipment delay rate • ASN submission compliance • Invoice matching success rate
Business Impact	• Inventory availability at dark stores • Procurement cost per purchase order • Vendor satisfaction score • Procurement team productivity • Stockout rate

Q If you had to choose one North Star Metric, what would it be?

A I'd choose Procurement Cycle Time, measured as the average time from purchase order creation to goods being received at the warehouse. Reducing this directly improves procurement efficiency, inventory availability, vendor collaboration, and ultimately reduces stockouts for customers.

Q Sounds good. We can conclude the discussion here.

WHAT WORKED · WHAT WAS MISSING

- ✓ **Three sharp clarifying questions.** Who uses it, replace or integrate, and which goal. "A layer over SAP, not a new ERP" shaped everything after it.
- ✓ **Stated the plan** and got a yes before starting.
- ✓ **Chose two of four users, with a reason:** procurement managers and vendors drive availability.
- ✓ **Ordered features by dependency.** Scorecards need data, so they come last. That is real product sense.
- ✓ **An explicit "leave out" list.** Saying what you won't build is half of prioritising.
- ✗ **Named RICE but scored nothing.** No reach, no impact numbers. Say "I'll order these by dependency" instead, or use the numbers.
- ✗ **Treated all vendors as one.** A national brand's key-account team on EDI and a local bakery taking orders by email need different products. Most vendors by count are the small ones.
- ✗ **Never sized the pain.** How many orders to suppliers go wrong, and what does a missing item cost a dark store?
- ✗ **Twenty metrics, no guardrail,** and a North Star that can improve while shelves stay empty.

Scoping kit check: Confirm the thing ✓ · Goal ✓ · Who ✓ · Today and the trigger ✓ (**replace or integrate**) · Edges: how many vendors, which cities ✗ **skipped**. Skipping the edges is why the vendor mix never came up.

CHECK · "NORTH STAR: PROCUREMENT CYCLE TIME" **OVERSTATED**

A faster purchase order (PO: the buyer's formal order to a supplier) helps little if the truck arrives with 70% of what was ordered. Quick commerce loses a sale when an item is missing at the one store near the customer. The metric closest to that is **OTIF**: the share of ordered lines delivered **on time and in full**, backed by the item fill rate at the store. Cycle time is a good input metric that feeds it.

CHECK · "EDI HAS LOWER REACH, AS NOT ALL VENDORS USE IT" **DIRECTIONALLY RIGHT**

EDI (electronic data interchange) is a standard format in which two companies' computers exchange orders and invoices directly. Large FMCG brands use it with modern retailers; local dairies and bakeries mostly don't. So EDI covers few vendors but a large share of purchase value. Measure reach in rupees of purchases, not number of vendors, and EDI jumps up the list.

STRONGER ANSWER



Figure 2.9 · A design answer in six moves. Moves ② and ④ are what the original skipped: two very different vendors, and a number for the pain. *Author's summary.*

A ① CLARIFY

Two more questions. How many vendors, and how concentrated are they? And what is the cost of today's mess: stock missing at dark stores, or staff time? I'll assume one country and all our cities.

Q About 3,000 vendors. The top 150 brands are 70% of purchase value; most of the rest use email or WhatsApp. The worry is stockouts: about 6% of items customers search for are missing at their store.

A ② **USERS**

So there are two vendors to design for. **Priya** runs key accounts at a national snacks brand: EDI, dozens of POs a day, and her pain is orders that fail silently. **Ramesh** runs a bakery supplying 60 dark stores: he reads POs in email at night and confirms on the phone, and waits 45 days to be paid. Our buyer **Neha** chases both.

Why this move: one product for Priya and Ramesh would suit neither. Naming them keeps every later choice honest.

A ③ **WORKFLOW**

The chain is: PO sent → vendor confirms quantity → truck dispatched with an advance shipping notice → store receives and counts → invoice matched → payment. Items go missing at two points: the vendor can't supply the full quantity and nobody hears until the truck arrives, and EDI messages fail without an alert.

A ④ **SIZE IT**

Illustrative: 30,000 PO lines a day × 8% arriving short or late = 2,400 lines. If each leaves an item missing in one store for a day, that's 2,400 store-days of lost sales, every day. Knowing a shortfall even 12 hours earlier lets Neha order from a second vendor.

Why this move: it turns "procurement is inefficient" into a number the fix can be judged against.

A ⑤ **PRIORITISE**

Using impact against effort (Figure 2.10). **Quick wins:** a mobile PO inbox where Ramesh confirms or changes quantities in one tap, and failure alerts on EDI orders for Priya. **Big bet next:** shipment tracking. **Later:** scorecards, once there are three months of data. **Not now:** AI forecasting and contract management.

A ⑥ **MEASURE**

Goal (North Star): PO lines delivered on time and in full (OTIF), from 92% to 96% in two quarters. **Inputs:** POs confirmed within 2 hours; share of POs on the platform; EDI error rate. **Safety check (guardrail):** active vendors must not fall, and store stockouts must not rise while we migrate. **Trade-off:** a portal adds work for small vendors, so email and WhatsApp stay as a fallback in the first version.

FRAMEWORK · IMPACT-EFFORT AND MOSCOW

1 · The situation. The candidate had ten sensible ideas and one team. A good list isn't a plan; a plan says what comes first and what never comes.

2 · An everyday example. Guests arrive in an hour and the flat is a mess. Clearing the living-room sofa takes 5 minutes and everyone will see it. Cleaning the bathroom takes 20 minutes and matters a lot. Repainting the scuffed wall takes a weekend and nobody will notice tonight. You sort by *how much it matters* against *how long it takes*, and you say out loud that the wall is not happening.

Part	Guests in an hour	Vendor platform
Quick win	Clear the sofa	One-tap PO confirmation on a phone
Big bet	Clean the bathroom	Live shipment tracking
Fill-in	Light a candle	Vendor scorecards (once data exists)
Money pit	Repaint the wall	Contract management, today
Won't (this time)	"Not tonight"	AI demand forecasting in the first version

3 · The framework. **Impact-Effort** places each idea on a 2×2: how much it moves the goal against how much work it takes. Start top-left (quick wins), plan the top-right (big bets), fit in the bottom-left, and avoid the bottom-right. **MoSCoW** then sorts a release into Must have, Should have, Could have and Won't have (this time). The Won't list is the useful part: it tells engineers, vendors and your boss what not to expect.




Figure 2.10 · Ten ideas, sorted twice. Left: impact on on-time-in-full deliveries against effort. Right: the same ideas as a release plan. Scorecards look like a quick win but can't work until the platform has data, so the 2x2 alone would have put them too early. Positions are illustrative judgments for this case.

4 · When it misleads. Both axes are guesses, and teams guess in their own favour: everything drifts into “quick wins”. The 2x2 also ignores order. You can't cook the curry before buying the onions, however quick the curry is; vendor scorecards are the onions-first trap here. MoSCoW breaks when 80% of the list becomes “Must”. A rule that helps: Musts may use at most 60% of the team's time.

5 · In the interview. “I'll place these on impact against effort, then call out one dependency: scorecards need data, so they wait. And here is what I'm deliberately not building in version one.”

CASE 2.2 · IMPROVE A FOOD DELIVERY APP

Product improvement

Medium IIM B Sigma PM Casebook 2026–27, pp. 126–127

Try it first. Think of the last time you opened a food app and closed it without ordering. At which step did you give up, and why? Build your answer from that moment.

THE QUESTION

We would like you to pick a food delivery app and focus on improving it. (Infoedge)

- A** I would like to choose Swiggy. It is a food delivery app which focuses on elevating the quality of life of Urban consumers by providing them convenience through food delivery. It has recently started its grocery delivery service too. It is considered to have a duopoly in the market and covid has led to changing habits of people which caused an overall increase in the size of the pie. This phenomena along with regulations has led to new players trying to enter the market.
- Q** That was comprehensive. Please proceed further.
- A** Is there a specific goal I should be able to achieve through the improvement features?
- Q** No, you are free to define the goal.
- A** Swiggy is a three way marketplace with 3 types of users - restaurants/grocery stores, end consumers and delivery partners. I would want to focus on end consumers, will that be okay?
- Q** Sure. Please proceed.
- A** It has verticals of food delivery, instamart and Swiggy Genie, are we specifically looking for one of these?
- Q** You can focus on one or more, it is solely your choice.
- A** Thank you for answering all the clarifying questions. I would like to define the scope of the problem before I proceed. I would like to focus on customer retention metrics and revenue. Since it is a popular name in the market with increased demand for delivery, customer acquisition doesn't seem to be a problem anymore. While defining improvements, I would also like to focus on enhancements which would serve as a competitive advantage. Will that be okay?
- Q** Sure. Go for it
- A** Since we have defined the goal, I would like to define user personas with their needs, goals and pain points. I would also like to focus on how and when do they really use our product. This is important to create value according to the context.
- Q** Sounds fair.

A User Personas can be defined as follows:

Families at Home: These are families with grandparents, parents, and kids – usually around 4 to 6 people. They usually cook their own meals and rarely order food, except for special occasions or when they want a break once a week. They like food that's good for both kids and older family members. They can afford to spend a bit more for high-quality and healthy food.

College Students: These are students who live away from home in college hostels. They get a fixed amount of money each month. The hostel food isn't the best, so they often try different food options. They're influenced by their friends and tend to eat at cheap stalls near the college to save money and hang out with friends.

Busy Professionals: These are people with busy jobs and big goals. They either don't like cooking or don't have time for it. They earn good salaries and don't have many responsibilities. They often order food about three-four times a week and go out to eat on weekends. They rely on recommendations from friends and colleagues for food choices.

Q Great! Please proceed.

A According to the context I would like to prioritize the user according to customer life time value. Households don't frequently order food, but their orders are usually larger. This presents a chance to offer them additional grocery items, increasing their long-term value. On the other hand, working individuals tend to order food frequently and prefer the convenience of not going to local stores for daily essentials. Given these insights into user preferences and behaviors, I would like to prioritize households and working individuals as the target audience. Is that fine?

Q Sure. Please proceed

The original shows three user-journey diagrams here; it is redrawn as Figure 2.11, labelled "redrawn from the original".

A Here are some challenges users face and certain beliefs they currently hold, which are affecting their optimal use of our product. Our aim is to transform these existing beliefs into more favorable ones:

- Perception of Grocery Quality and Cost: Some users consider buying groceries through Swiggy to be pricey and worry that the freshness might not be up to the mark.
- Discovering New Places: Users feel that Swiggy lacks information about recently opened restaurants in their vicinity, especially those offering the types of dishes they usually enjoy.
- Hesitation due to Hunger: Users are reluctant to place orders when they're very hungry, as browsing through options and delivery takes a significant amount of time.
- Emotional Connection with Food: The platform is perceived as lacking the emotional connection that fancy restaurants provide, appearing more transactional in nature.
- Complex Ordering for Households: In households, the process of taking everyone's food preferences becomes a hassle, leading to longer ordering times. This problem escalates when hosting parties or having friends over.
- Reliability of Ratings: Users find it hard to trust restaurant ratings since nearly all establishments seem to have high ratings, making the authenticity of these ratings questionable.

The casebook repeats an earlier interviewer line here ("You can focus on one or more, it is solely your choice."); the repeat is removed.

A First, I would like to define solution to these problems and then prioritize using RICE framework

- Sampling and Grocery Incentives: To engage households who order from restaurants, we could offer a complimentary Instamart sample. Coupled with coupons, this could encourage them to explore Swiggy's grocery offerings. After delivery, a push notification can inquire about their experience with the sample and nudge them to visit Instamart for groceries. This might require delivery partners to pick up samples during downtimes, with extra compensation.
- Restaurant Collaborations and Emotional Content: Collaborating with new or recently opened restaurants presents an opportunity. Short videos could provide insights into the food, chefs' stories, and even delivery partner experiences. Such content, along with brief restaurant descriptions in menus, fosters an emotional connection. Users could also share this content with their social circles.
- Voice Ordering: Integration with virtual assistants like Alexa, Google Assistant, or Siri for hands-free ordering. This would be extremely helpful for adding multiple groceries to instamart cart in an instant.
- Group Ordering and Scheduled Deliveries: Introducing a group order button on menu pages, generating a time-limited link for shared orders, adds a sense of urgency. Implementing scheduling features, although complex, would enhance convenience. Balancing item availability, restaurant timing, and delivery partners' schedules is crucial.
- Variable Rewards for Engagement: Utilizing a variable reward concept post-order could maintain user engagement. Random coupon or scratch card rewards provide personalized discounts and offers, deviating from generic promotions. In the next step I would like to prioritize these features. Prioritization using the RICE framework is a logical step. Ensuring win-win-win situations, where consumers, restaurants, and delivery partners all benefit, aligns with Swiggy's approach. Each proposed project or feature will be assessed based on its potential positive impact on all stakeholders.

Q Sure. Sounds good. Please proceed with the roadmap.




Figure 2.11 · The candidate's three user journeys, redrawn from the original. Each row is one reason to open the app, broken into the steps the user takes. The steps are right; what is missing is how the user feels at each one, which is where the pain points should come from. *Redrawn from the original, IIM B Sigma PM Casebook 2026–27, pp. 126–127.*

WHAT WORKED · WHAT WAS MISSING

- ✓ **Asked for the goal**, then picked one and gave a reason: acquisition isn't Swiggy's problem, so retention and revenue.
- ✓ **Chose one side** of the three-sided marketplace and checked it with the interviewer.
- ✓ **Personas that differ** in how often and how much they order.
- ✓ **Used journeys** for three reasons to open the app, and framed pains as beliefs to change.
- ✗ **A minute of company history** nobody asked for. Confirm the product in one line.
- ✗ **The priority logic contradicts itself** (see the Check below).
- ✗ **Pain points float free of the journey.** None is tied to a step, and none is called the worst.
- ✗ **Five ideas, no ranking.** The case ends before the promised RICE. Voice ordering through Alexa solves a problem no persona had.
- ✗ **"Retention and revenue" never became a number.**

Scoping kit check: Confirm the thing ✓ (**too long**) · Goal ✓ · Who ✓ · Today and the trigger: what's changed, where are we losing? ✗ **skipped** · Edges ✓ (**which vertical**). Without a trigger the goal stayed vague.

CHECK · "HOUSEHOLDS ORDER RARELY BUT SPEND MORE, SO PRIORITISE THEM BY LIFETIME VALUE" **DOESN'T HOLD UP**

Lifetime value is orders × margin × how long they stay, so frequency usually beats basket size. Illustrative, at 8% contribution: a household ordering twice a month at ₹900 is worth $2 \times ₹900 \times 8\% = ₹144$ a month. A professional ordering 12 times at ₹350 is worth $12 \times ₹350 \times 8\% = ₹336$ a month, more than twice as much. The households may still be a good bet for Instamart, but not on this logic.

CHECK · "SWIGGY IS CONSIDERED TO HAVE A DUOPOLY" **HOLDS UP**

Zomato and Swiggy handle the large majority of India's app-based food orders, roughly 55:45 between them by order value in 2025–26. ONDC apps such as magicpin and newer entrants are growing from a small base. The useful follow-on is the gap: Zomato food grew 20% in Apr–Jun 2026, Swiggy 17%.

STRONGER ANSWER



Figure 2.12 · An improvement answer in six moves. The map in move ③ exists to find move ④: the one step where the user feels worst. *Author's summary.*

A ① CLARIFY

I'll take Swiggy's food app in India, for customers. Swiggy is the close number two and grew slower than Zomato last quarter, so I'd aim at **orders per monthly user** from people who already use it. Has anything changed recently that I should know about?

Q Nothing specific. Your call on the rest.

A ② ONE USER

Then the user is **Arjun**, 29, who works late and orders dinner three or four times a week. He is Swiggy's most valuable customer and the easiest to lose: Zomato is one tap away and has the same restaurants.

A ③ MAP

His weeknight journey: feel hungry → decide what to eat → pick a restaurant → pay → wait → eat → rate. I'll mark how he feels at each step (Figure 2.13). Two low points: deciding (12 minutes of scrolling through 40 look-alike options at 8:45 pm) and waiting (the promised 31 minutes becomes 45).

A ④ WORST STEP

Illustrative: of 100 evening sessions, 55 end in an order. Of the 45 that don't, most leave while browsing, before a cart exists. So deciding is the worst step, and waiting is second.

Why this move: the map is only useful once it names one step to fix first.

A ⑤ SOLVE

"Tonight's three": between 7 and 11 pm the app opens on three one-tap picks: his most-reordered dish, one new dish close to his taste, and the fastest reliable option near him. **An honest ETA**: a range (30–38 min) and an on-time badge for restaurants that keep their promise.

A ⑥ MEASURE

Goal (North Star): orders per monthly user among weeknight users, from 3.4 to 3.7. **Inputs**: time from app open to payment (12 → 6 minutes); share of evening sessions ending in an order. **Safety check (guardrail)**: order value per order, and cancellations. **Trade-off**: three picks can bury small restaurants that pay for ads, so one of the three is always a discovery slot.

FRAMEWORK · THE CUSTOMER JOURNEY MAP

1 · The situation. The candidate drew the steps of three journeys, then listed pain points separately. Nobody could tell which step hurt most, so nothing got fixed first.

2 · An everyday example. Think of an overnight train from Pune to Delhi. Booking: anxious (waitlisted). Reaching the station: rushed. Finding the coach: confused. The night: fine. Arrival at 6 am: tired. If you ran Indian Railways and could fix one thing, better food on board would not be it. The low point is the waitlist.

Row of the map	Train to Delhi	Arjun's weeknight dinner
Stages	Book → reach station → find coach → travel → arrive	Hungry → decide → pick → pay → wait → eat → rate
Doing	Refreshes the waitlist status	Scrolls 40 options
Feeling	Lowest at booking	Lowest at deciding, then waiting
Pain point	"Will I even get a berth?"	"Everything looks the same, it's 8:45"
Opportunity	Confirmed-berth prediction	Tonight's three, one tap each

3 · The framework. A **customer journey map** lays out the stages a user goes through to get a job done, and for each stage what they do, think and feel. The feelings, drawn as a curve, show the low points. Pain points and opportunities are written under the stage where they occur, so every idea is tied to a moment.




Figure 2.13 · Arjun's weeknight dinner as a journey map. Read the curve first: the two red dips, deciding and waiting, are where he quits or gets angry. Ideas sit under the step they fix, so each one has a reason to exist. *Illustrative persona and journey.*

4 · When it misleads. A map drawn at the team's desk shows the team's guesses. Plan a wedding from the host's desk and you will map the ceremony beautifully and never see the guests' two-hour parking queue. Build the map from real evidence: session recordings, where users drop off, support tickets, five user interviews. And map the unhappy path too: the late order, the wrong item.

5 · In the interview. "Let me walk one weeknight order step by step and mark where Arjun feels worst. I'll fix that step first."

CASE 2.3 · METRICS FOR BIGBASKET'S APP

Metrics Medium IIM A Product Mindset 2026–27, pp. 291–292

Try it first. Your family buys groceries online. Write the one number that would tell BigBasket it is winning your household, then read on.

THE QUESTION

You are the Digital Product Manager of Big Basket. What would be the metrics you would like to follow?

- A** I'd first like to confirm my understanding of BigBasket. BigBasket is an online grocery platform where users can order groceries, fruits & vegetables, packaged foods, household essentials, and personal care products through its mobile app or website.
- Q** Yes, you may move forward.
- A** Before defining metrics, could you share the business goal we'd like to measure?
- Q** Assume that the goal of the business is to increase the value of orders placed via the BigBasket mobile app.
- A** Understood. I'll first break down the business goal into measurable drivers and then map the user journey to identify the right metrics. Should I cover both the website and the mobile app?
- Q** You can define the metrics for only the mobile app for now.
- A** Sure. The business goal can be represented as:
- Total Order Value = Number of Orders × Average Order Value (AOV)
 - Number of Orders = Visitors × Add-to-Cart Rate × Purchase Conversion Rate
 - Visitors = New Visitors + Returning Visitors
- I'll now map the user journey. The mobile app journey would broadly be:
- Install the app
 - Open the app
 - Browse products via home page, categories or search
 - Add products to cart or save them to a list
 - Log in / Sign up (if required)
 - Review cart, choose address & delivery slot, apply coupons or opt for BBStar
 - Complete payment
 - Receive order confirmation
 - Track, modify or cancel the order through the app
 - Receive delivery and rate the experience
- Based on this, I'd divide the metrics into five buckets. Awareness Measures whether users are discovering and installing the app.
- Number of app installs per month
 - % installs from organic sources (Search, Website)
 - % installs from paid campaigns (Ads, Influencers)
 - App Store ranking in Shopping/Grocery
 - Click-to-install conversion rate
 - Monthly uninstall rate
 - Average App Store rating
- Engagement Measures how actively users explore products before purchasing.
- Bounce rate by product/category page
 - Average session duration
 - Average products viewed per session
 - Average quantity added to cart
 - Number of custom lists/wishlists created
 - % users using in-app search
 - % logged-in users
 - % new registrations
 - Daily Active Users (DAU)
 - Monthly Active Users (MAU)
 - Monthly value of returned orders
- Conversion Measures how efficiently browsing converts into completed purchases.
- Average number of items per order
 - Average basket value
 - Cart abandonment rate
 - Order cancellation rate
 - Average Order Value (AOV)
 - % users accepting product recommendations during checkout
 - % users applying coupons/vouchers

- Average voucher value redeemed

Retention Measures whether users continue purchasing through the app. % users subscribing to BBStar

- Repeat purchase rate
- Purchase frequency (orders/user/month)
- Average time between two purchases

Referral Measures growth driven by existing customers.

- % users sharing referral codes
- Referral conversion rate
- Revenue generated through referrals
- Net Promoter Score (NPS)

Q This is great. Thank you!

WHAT WORKED · WHAT WAS MISSING

- ✓ **Confirmed the product in one line** and asked for the goal before listing a single metric.
- ✓ **Wrote the goal as a formula:** order value = orders × AOV, and orders broken down further. That is a driver tree.
- ✓ **Scoped the platform** (app only) with one question.
- ✓ **Walked the journey first**, so the metrics follow what the user does.
- ✗ **About 40 metrics, none chosen.** No North Star, no guardrail, no target.
- ✗ **The buckets are AARRR without Revenue**, although order value was the stated goal.
- ✗ **The formula forgets frequency**, the biggest lever in grocery (see Check).
- ✗ **BB Now (10–20 minutes) and scheduled slots** were never split, though their baskets differ a lot.

Scoping kit check: Confirm the thing ✓ · Goal ✓ · Who: new or repeat buyers? ✗ **skipped** · Today and the trigger ✗ **skipped** · Edges ✓ (app only). Metrics add-on, “what decision will these metrics drive?”: ✗ **skipped**, which is how a list of 40 happens.

CHECK · “ORDERS = VISITORS × ADD-TO-CART RATE × PURCHASE CONVERSION” **DIRECTIONALLY RIGHT**

True for one visit. But groceries are a habit: a household that orders every week and four strangers who order once look the same in this formula. For a month, a better split is **orders = monthly buyers × orders per buyer**. It shows at once whether growth is coming from new households or from existing ones ordering more often.

CHECK · “DAILY ACTIVE USERS AS AN ENGAGEMENT METRIC” **OVERSTATED**

Families buy groceries three to six times a month. A daily-use ratio would call a healthy weekly habit “weak engagement” and push the team towards notifications nobody wants. Match the metric to the natural rhythm of the job: here, orders per buyer per month.

STRONGER ANSWER



Figure 2.14 · A metrics answer in six moves. The list of metrics shrinks from 40 to about eight, and each one earns its place. *Author's summary.*

A ① CLARIFY

Over what time frame, and should I include both BB Now and scheduled delivery? I'll assume the goal is order value from the app over the next two quarters, from both new and existing buyers.

Q Both services, two quarters. Go ahead.

A ② FORMULA

Monthly app order value = monthly buyers × orders per buyer × AOV. Illustrative: 40 lakh × 3.5 × ₹750 = ₹1,050 Cr a month. Lifting any one factor by 10% adds ₹105 Cr. For households that already buy, frequency is the cheapest to move.

Why this move: the formula shows which lever is worth the most before any metric is named.

A ③ STAGES

I'll follow the buyer through five stages (Figure 2.15): they find the app, place a first order, come back, spend more, and tell others.

A ④ ONE EACH

Acquisition: cost per first order. **Activation:** share of new installers who order within 7 days. **Retention:** share of first-time buyers who order again within 30 days. **Revenue:** orders per buyer per month, and AOV. **Referral:** invites accepted per 100 buyers.

A ⑤ NORTH STAR

Goal (North Star): monthly order value from repeat buyers (households with two or more orders in the month). It grows only if people come back and buy more, which is what the business goal needs.

A 6 GUARDRAILS

Safety checks (guardrails): on-time delivery %; missing or substituted items per 100 orders; contribution margin per order, so discounts don't buy the growth. **Trade-off:** pushing BB Now raises frequency but shrinks baskets, so I'd watch the mix. **Next:** a weekly review of the five stage metrics, split by service.


FRAMEWORK · AARRR, THE CUSTOMER FUNNEL

1 · The situation. The candidate grouped 40 metrics into buckets. The buckets were almost right; what was missing was one number per bucket and the reason for it.

2 · An everyday example. A new gym opens in your colony. Some people see the poster. A few come for a free trial. Some enjoy the first workout. Fewer are still coming in week six. Those pay for a six-month plan. A couple bring a friend. If the gym is failing, it matters a lot which of those steps is leaking.

Stage	The gym	BigBasket app	One metric
Acquisition	Sees the poster, visits	Installs and opens the app	Cost per first order
Activation	First workout feels good	First order arrives fresh, on time	Order within 7 days of install
Retention	Still coming in week six	Orders again this month	Second order within 30 days
Revenue	Buys the six-month plan	Bigger, more frequent baskets; BB Star	Orders per buyer × AOV
Referral	Brings a friend	Shares a referral code	Invites accepted per 100 buyers

3 · The framework. AARRR (Dave McClure's "pirate metrics", 2007) splits a customer's life into five stages: Acquisition, Activation, Retention, Revenue, Referral. Measure the conversion from each stage to the next, find the biggest leak, and work there. **Activation** means the first moment the user gets the promised value, here a first order that arrives complete.



Illustrative monthly numbers for BigBasket's app.

Figure 2.15 · One month of new users, stage by stage. The funnel narrows fastest between install and first order, so that is where the next quarter's effort goes. Retention is where the money is. *Illustrative numbers. Framework: Dave McClure, "Startup Metrics for Pirates", 2007.*

4 · When it misleads. The funnel invites a list of metrics per stage, which is how this candidate reached 40. It also draws acquisition at the top, as if it mattered most. A gym that spends on posters while members stop coming in week two is filling a leaking bucket. In grocery the order of attention is often retention, activation, revenue, and acquisition last.

5 · In the interview. "I'll walk the five stages and pick one metric for each. Then I'll choose the North Star from the stage that moves order value most, which for groceries is repeat buying."

CASE 2.4 · SWIGGY'S PRODUCT METRICS BEFORE ITS IPO

Metrics Medium IIM B Sigma PM Casebook 2026–27, p. 153

Try it first. You are about to buy Swiggy shares. Its listed rival, Zomato, already makes a profit. Which three numbers would you ask Swiggy for before you invest?

THE QUESTION

You are the CPO of Swiggy, what key product metrics would you monitor in preparation for an IPO?

A To ensure a successful IPO, I would focus on product metrics that provide valuable insights into user engagement, satisfaction, and overall platform performance, these factors would help convince investors. Does this sound reasonable, should I proceed and list the metrics?

Q Yes, please proceed. Also, include the rationale behind selecting each metric, particularly in the context of the upcoming IPO.

A 1. Monthly Active Users (MAUs): MAU indicates overall user engagement and growth. It is important for IPO as investors look for strong user growth and engagement as a sign of a successful and scalable product.

2. Daily Active Users (DAUs): DAU measures daily user engagement and platform stickiness. It becomes important for IPO because a high DAU to MAU ratio can signal strong user retention and satisfaction, a key factors that can help attract investors.

3. Customer Retention Rate: It shows how well the platform can keep its users over time. A strong retention rates can help demonstrate long-term user loyalty and the product's potential for sustained revenue.
4. Order Volume and Growth: Order volume tracks the total number of orders placed, reflecting market penetration and user activity. An increasing order volume showcases the platform's growth trajectory and can help appeal to investors looking for scalable businesses.
5. Average Order Value (AOV): AOV helps measure average spending per order. A high AOV contributes to overall revenue and will help highlight the platform's potential for profitability.
6. Average Delivery Time: It helps track the efficiency of the delivery process. Hence, critical for IPO as efficient delivery operations enhance user satisfaction and operational excellence, key considerations for investors.
7. Feature Usage and Engagement: Feature-wise insights into user engagement. Important for IPO because identifying popular features helps enhance user experience and engagement, attracting investors.
8. Customer Satisfaction (Net Promoter Score - NPS): NPS measures likelihood of users recommending the platform. A high NPS indicates strong user satisfaction and loyalty, vital for growth. Therefore, vital for IPO
9. Churn Rate: It measures the percentage of users who stop using the platform over time. Important for IPO because managing churn is essential for long-term user base stability and revenue predictability, factors closely monitored by investors.
10. Operational Efficiency Metrics: Assessing how efficiently the platform operates, including fulfillment and staffing costs is again critical. Important for IPO as investors value companies that manage costs effectively while scaling, indicating a sustainable business model.

Q It was a comprehensive list. How would you use these metrics to prepare for the IPO?

A I would closely monitor these metrics and utilize them to:

1. Demonstrate Growth: Use positive trends in MAUs, DAUs, and order volume to showcase the platform's growth and user engagement.
2. Highlight Efficiency: I will showcase improvements in delivery times and operational efficiency to underline our ability to scale effectively.
3. Emphasize User Satisfaction: I will use high NPS and retention rates to reflect strong user satisfaction and loyalty.
4. Provide Financial Insights: Finally, linking metrics like AOV and feature usage to revenue potential and profitability should be insightful.

Q This was comprehensive solution.

A Thank you.

WHAT WORKED · WHAT WAS MISSING

- ✓ **Checked the approach** and gave a reason for each metric, as asked.
- ✓ **Covered growth, engagement, satisfaction and efficiency**, so nothing major is left out entirely.
- ✓ **The last turn tells a story**: growth, efficiency, satisfaction, then financials. An IPO is a story told with numbers.
- ✗ **Every metric would fit any consumer app**. Nothing about food, Instamart, riders or dark stores.
- ✗ **No money per order**. Contribution margin, the share of order value Swiggy keeps as revenue, and cash burn were what investors asked about most in 2024.
- ✗ **No split by business**. Food made money; Instamart lost a lot. One blended number hides both.
- ✗ **No comparison with Zomato**, the number every investor would hold up against Swiggy's.

Scoping kit check: Confirm the thing ✗ · Goal: what does a “successful IPO” mean here? ✗ · Who: which investors, which businesses? ✗ · Today and the trigger: Zomato already listed and profitable ✗ · Edges ✗. Five skips in a row are why the list came out generic.

CHECK · “A HIGH DAU-TO-MAU RATIO SIGNALS STRONG RETENTION” **OVERSTATED**

People order food three or four times a month, not daily. A food app with a low daily-to-monthly ratio can have excellent customers. The retention evidence investors trust is **cohort retention**: of the people who first ordered in January, what share still order in June, and how often.

CHECK · “A HIGH AOV HIGHLIGHTS THE POTENTIAL FOR PROFITABILITY” **DIRECTIONALLY RIGHT**

A bigger basket spreads the fixed cost of the ride. But a basket made bigger by a ₹100 discount can lose more money, not less. Profitability shows up in contribution per order, and AOV is only one input to it.

STRONGER ANSWER



Figure 2.16 · An IPO-metrics answer in six moves. Moves ④ and ⑤ answer the question every investor brings: does each order, and each customer, make money? *Author's summary.*

A ① CLARIFY

Which businesses are in scope, food and Instamart? And what is the main doubt investors have about us?

Q Both. They will compare us with Zomato, which is already profitable.

A ② THE QUESTION

Then the story must answer one doubt: can Swiggy grow and make money on each order in both businesses? I'll show growth, then money per order, then what a customer is worth.

A ③ GROWTH

Order value = monthly transacting users × orders per user × AOV, shown separately for food and Instamart. For users, cohort curves: of January's new customers, how many still order in June.

A ④ PER ORDER

Contribution margin per order in each business, with the trend. Illustrative: food ₹35 and rising; Instamart from – ₹40 towards zero as dark stores mature. Next to it: adjusted EBITDA as a % of order value, and the share of dark stores older than a year.

Why this move: for a loss-making company, the path of money per order is the whole case.

A ⑤ PER CUSTOMER

Illustrative: it costs ₹400 to win a new customer. Each then brings ₹25 per order × 3 orders a month, on a falling retention curve, so they pay back in month 12 (Figure 2.17). Over two years that is about ₹640 against ₹400. Customers won with heavy discounts never pay back, so I'd show payback by channel.

A ⑥ GUARDRAILS

Goal (North Star) of the story: contribution per order, positive and rising in both businesses, while order value grows 20%+. **Safety checks (guardrails):** on-time %, rider earnings per hour, and complaints, so margins aren't bought by squeezing riders or quality. **Trade-off:** cutting discounts lifts margin but slows growth right before listing.

FRAMEWORK · UNIT ECONOMICS: CONTRIBUTION MARGIN, LTV/CAC AND PAYBACK

1 · The situation. Investors buying a loss-making company want to know whether losses shrink as it grows. Only per-order and per-customer numbers answer that.

2 · An everyday example. A chai stall sells a cup for ₹15. Tea, milk and sugar cost ₹5, the paper cup ₹1, so each cup contributes ₹9. Rent and a helper cost ₹18,000 a month, so the stall breaks even at 2,000 cups, about 67 a day. To win regulars, the owner gives a free first cup (₹6 of cost). A regular who then drinks 20 cups a month earns that back in a day; a passer-by who never returns does not.

Idea	Chai stall	Swiggy food
Contribution per unit	₹15 – ₹6 = ₹9 a cup	~₹25–35 an order, after rider, discounts, payments
Fixed costs	Rent + helper, ₹18,000 a month	Tech, staff, brand marketing
CAC (cost to acquire a customer)	A free first cup, ₹6	~₹400 of marketing per new customer (illustrative)
LTV (lifetime value)	Cups a regular buys × ₹9	Orders over their life × contribution per order
Payback	One day	About 12 months

3 · The framework. **Contribution margin** is money per unit after the costs that unit causes. **CAC** is marketing spend ÷ new customers won. **LTV** adds up the contribution a customer brings over their life, allowing for the fact that many stop ordering. A common rule of thumb asks for LTV at least three times CAC. **Payback** is the month in which a customer's cumulative contribution passes their CAC; the shorter, the less cash the company needs to grow.




Figure 2.17 · When does a new customer pay for themselves? Green: a customer who found the app on their own pays back the ₹400 in a year. Amber: a customer won with a large discount orders less often and stops sooner, and never covers the ₹400. Same average CAC, very different businesses. *Illustrative model: ₹25 contribution per order, 3 orders a month, retention falling to about 30% by month 12.*

4 · When it misleads. LTV is a forecast, and a forecast of five years of loyalty is mostly hope. Averages also hide cohorts: the chai stall that counts the Diwali-week crowd as regulars will think its free cups are working. And companies define “contribution” differently (some leave out discounts or customer support), so compare like with like before comparing Swiggy with Zomato.

5 · In the interview. “Growth only helps if each order makes money and each customer pays back their acquisition cost. So I’ll show contribution per order by business, and payback by acquisition channel.”

CASE 2.5 · LAUNCHING TABLE RESERVATIONS ON SWIGGY Go-to-market Medium IIM A Product Mindset 2026–27, pp. 280–282

Try it first. Who in your circle books a restaurant table in advance, and for what occasion? Name one group of people and one reason they would book through Swiggy rather than call the restaurant.

THE QUESTION

Swiggy has decided to enter the table reservation category in restaurants. Formulate a launch strategy for the same.

A Before structuring the GTM, a few clarifying questions:

1. Build approach - Should this be a new standalone app or a feature within the existing Swiggy app?
2. Stage of development - Is the product defined, or are we still in the conceptualisation phase?
3. Target geography - Which cities are we launching in first?
4. Goal - Is the primary objective user acquisition, restaurant partner growth, or revenue?

Q The product is still in the idea conceptualisation phase - I was hoping you could help with that too. It should be a feature within the existing Swiggy app.

A Good. Building within the existing Swiggy app is the right call. It avoids the cold-start problem of a new app - Swiggy already has 90M+ users and strong restaurant relationships. We can distribute instantly without any install friction. Before moving to the GTM, let me map the stakeholders and their pain points - this will ground the feature design and the launch strategy in real user needs:

Stakeholder	Pain Points	What the Feature Solves
Customers	- Don't know wait times, cuisines, or table availability before arriving - Dislike queuing - especially for planned occasions - Want parking and timing information upfront	- Browse restaurants, check real-time availability, and book a table in advance - Pre-order dishes during booking to reduce wait after arrival - Arrive with confidence - no queue, no uncertainty
Restaurants	- No-shows and late arrivals block tables and cost revenue - Unpredictable demand makes staffing and prep difficult - Waiting customers are lost if competitors are nearby	- Estimated arrival times help kitchens prepare orders ahead of time - Buffer time policy protects against table-blocking without penalising punctual guests - Better demand visibility enables staffing and inventory planning
Swiggy	- Currently absent from the dine-in occasion - losing engagement to Zomato - Existing user base not monetised beyond food delivery	- Expands Swiggy's presence into the dine-in journey - Increases DAU and session frequency through a new use case - New revenue stream via booking fees or premium restaurant listings

Q But doesn't Zomato already provide table reservation features?

A Yes - and that's an important competitive signal. But the key differentiator here is pre-ordering. Zomato's reservation feature lets users book a table, but the order is still placed at the restaurant. Swiggy's feature can allow customers to pre-order dishes during the booking flow, so the kitchen starts preparing before they arrive. This reduces wait time significantly and turns the reservation into a full-service, seamless dine-in experience - something Zomato doesn't currently offer at scale. This also addresses the restaurant's biggest fear about reservations: blocking a table for a customer who might not show up or arrives late. To manage this, we can introduce a configurable buffer window (e.g. 15 minutes) - the restaurant holds the table for the duration of the buffer. If the customer arrives after the buffer expires and the restaurant is full, they join the walk-in queue. This protects restaurant revenue without penalising punctual guests.

Q Considering this to be what the application does, how would you go about launching it?

A I'll structure the launch across three phases - Pre-Launch, Launch, and Post-Launch - covering Product & Pricing, Promotion, and Distribution for each:

Phase	Product & Pricing	Promotion	Distribution
Pre-Launch	- Launch as a feature within the existing Swiggy app - no new install required for existing users - No booking fee initially- free for customers and restaurants to lower adoption friction	- In-app banners and push notifications to existing Swiggy user base - Teaser campaign on Instagram and YouTube targeting urban millennials - Outreach to 200-300 premium restaurants in Mumbai, Bengaluru, and Delhi for early onboarding	- Leverage existing Swiggy app distribution- 90M+ users already installed - Restaurant partner sign-up via dedicated onboarding team and self-serve portal
Launch	- Introduce referral incentive: existing users get a discount voucher for every new user who makes their first reservation through a shared link - Exclusive launch offers -flat discounts or complimentary desserts- negotiated with partner restaurants for early adopters	- Smart nudges in the app: surface the reservation feature when users search for a restaurant (not just food delivery) - Social sharing of reservations to drive organic awareness - Targeted ads in dining-heavy localities and near premium restaurant clusters	- App Store and Google Play via existing Swiggy app update - Restaurant partners promoted as 'Swiggy Reserve' outlets -visible badge in search and discovery
Post-Launch	- Introduce convenience fee (₹20-₹50 per booking) once the feature has critical mass - Premium restaurant listing - paid placement for restaurants wanting higher visibility - Group booking and corporate dining packages	- Re-engagement campaigns for users who browsed but didn't book - Feature-led PR: 'X million tables reserved on Swiggy' milestone content - Partner with food bloggers and review platforms for discovery	- Expand to Tier 2 cities (Pune, Hyderabad, Ahmedabad) based on adoption data - Integrate with Google Maps and Apple Maps for restaurant discovery outside the Swiggy app - Explore corporate tie-ups for team dining and client entertainment use cases

A Success Metrics

I'd track performance across six dimensions - with reliability metrics being the most critical, since a feature that fails to hold reservations destroys trust faster than it builds it:

Stage	Metric	Why It Matters
Acquisition	# of reservations made; Cost per first booking; Conversion rate from feature page views	Measures whether the feature is discoverable and compelling enough to drive action
Engagement	# of reservations per active user per month; % of reservations with pre-orders; Bounce rate on reservation page	Pre-order attachment rate is the key differentiator vs Zomato - it validates the unique value proposition
Revenue	Booking fee revenue (post-launch); Premium listing revenue; ARPU from reservation users	Validates the path to monetisation once the feature has proven user and restaurant value
Supply Quality	# of restaurant partners onboarded; % of listed restaurants with real-time availability	A thin restaurant supply kills the feature - this tracks the health of the two-sided marketplace
Reliability	% of customers who arrived within the buffer window; No-show rate; Restaurant cancellation rate	Trust is the core product - if reservations don't hold, the feature fails regardless of acquisition numbers
Retention	% of users who make a second reservation within 30 days; Feature DAU/MAU ratio	Repeat booking is the signal that the product has created a new habit, not just a one-time trial

A The most important early signal will be the repeat booking rate within 30 days. A single reservation is a trial. A second reservation is a habit. If repeat bookings are low, it signals either a supply problem (not enough restaurants) or a reliability problem (reservations not being honoured) - both require different interventions.

Expansion Roadmap

Based on performance data from the initial metro launch, I'd plan expansion in three directions:

- Geographic - roll out to Tier 2 cities (Pune, Hyderabad, Ahmedabad, Chandigarh) where dining out is growing but table reservation tools are absent.
- Vertical - extend to cafes, cloud kitchens with dine-in sections, and corporate dining - broadening the addressable restaurant base.

- Monetisation - introduce a per-booking convenience fee (₹20-₹50) once NPS and reliability metrics confirm the feature is trusted. Layer in premium restaurant listings and corporate dining packages.

Q Sounds good. We can close the case now.

WHAT WORKED · WHAT WAS MISSING

- ✓ **Four clarifying questions:** build approach, stage, geography and goal.
- ✓ **Named the restaurant's real fear,** no-shows, and answered it with a buffer window.
- ✓ **Handled the "Zomato already does this" push** with a clear difference: pre-ordering.
- ✓ **A phased plan** across product, promotion and distribution, and one metric singled out: repeat bookings within 30 days.
- ✗ **Missed Swiggy's own asset.** Swiggy bought Dineout in 2022 and runs dine-out inside its app. The launch should start from those restaurants and deals.
- ✗ **The goal question got no answer,** and the candidate didn't push. Without a goal there is no way to choose between users, restaurants and revenue.
- ✗ **No target segment.** "Urban millennials" and "premium restaurants" describe everyone and no one.
- ✗ **Channels listed, none tested.** Instagram, YouTube, bloggers, maps and PR all at once is a budget, not a plan.

Scoping kit check (GTM add-ons): Objective ✓ **asked** ✗ **not pinned down** · Standalone or add-on ✓ · Target market ✓ **cities** ✗ **segment** · Competitors ✓ **when pushed** · Budget ✗ **skipped**.

CHECK · "SWIGGY ALREADY HAS 90M+ USERS, SO DISTRIBUTION IS INSTANT" **OVERSTATED**

A figure like that counts everyone who has ever installed or ordered. The people who actually order in a month number closer to 2 crore (2025-26). And a push notification about table booking reaches a delivery habit, not a dining-out one. Reach is real, but an in-app banner is a channel to test, not a guarantee.

CHECK · "MONETISE LATER WITH A ₹20-50 BOOKING FEE" **OVERSTATED**

Indian diners rarely pay to reserve a table that the restaurant would give them free on the phone. Dine-out apps earn mainly from restaurants: a commission when the bill is paid through the app (usually with a discount the restaurant funds) and paid listings. A booking fee fits only scarce slots, such as New Year's Eve.

STRONGER ANSWER



Figure 2.18 · A launch answer in six moves. The original jumped from features to a phased plan; moves ② and ③ decide who the plan is for. *Author's summary.*

A ① CLARIFY

Before the plan: is success more dine-out occasions on Swiggy, more restaurants, or revenue in year one? I'll assume the goal is to make dining out a habit on Swiggy, counted as bookings that were actually honoured. And is the budget large or modest?

Q That goal is right. Assume a modest budget.

A ② ASSETS

Swiggy already has Dineout: restaurants signed up, bill-payment deals and Swiggy One members. So this is reservations with pre-ordering added to an existing business, not a new category. That cuts the cost of getting restaurants on board.

Why this move: "what do we already own?" is the first GTM question, and here it changes the plan.

A ③ SEGMENT

Who books tables, and for what (Figure 2.19)? Friday friend groups, couples on date nights, families at Sunday lunch, office teams at weekday lunch. I'd target **office lunch groups in Bengaluru** first: they have an hour, they hate waiting for food, and pre-ordering fixes exactly that.

A ④ OFFER

Position: "For teams with an hour for lunch, Swiggy books the table and has the food ready when you sit down." For restaurants: pay only for honoured bookings, and a ₹100 per person deposit, refunded on arrival, to cut no-shows.

A ⑤ CHANNELS

List many channels, then test three cheap ones for three weeks: a banner at 11:30 am for users who order lunch to office addresses; tie-ups with office admin teams; a QR on restaurant tables that says "book lunch next time". Keep the one with the lowest cost per honoured booking.

A ⑥ MEASURE

Goal (North Star): honoured bookings per month. **Inputs:** share of bookings with a pre-order; restaurants with open lunch slots. **Safety check (guardrail):** no-show rate and restaurants leaving. **Trade-off:** deposits cut no-shows but also bookings, so test ₹0 against ₹100.

FRAMEWORK · STP, THE FIVE GTM QUESTIONS, AND BULLSEYE

1 · The situation. The candidate had a good product and a phased plan, but launched it at “urban millennials” through every channel at once. A launch needs one group to win first and one channel that works.

2 · An everyday example. Your aunt starts a home tiffin service in Koramangala. Students want cheap, bachelors want home-style food at the office, elderly couples want less oil. She picks working bachelors in three PG hostels, promises “home food at your desk by 1 pm”, and tries three ways to find them for two weeks: posters in the PGs, a WhatsApp group, and the PG cook. The WhatsApp group brings nine of every ten customers, so she stops the posters.

Step	Tiffin service	Swiggy reservations
Segment	Students, bachelors, elderly couples	Friend groups, couples, families, office teams
Target	Bachelors in three PGs	Office lunch groups, Bengaluru
Position	“Home food at your desk by 1 pm”	“Table booked, food ready when you sit”
Channel test	Posters vs WhatsApp vs the PG cook	Banner vs office admins vs table QR

3 · The framework. STP means Segment (split the market by need or situation), Target (choose the segment you can win first) and Position (one line on why you, for them). It answers the first of the **five GTM questions**: **Who** is it for? **What** do we offer them? **Why us, why now** (our assets)? **How** will they hear about it and buy it (channel and price)? **How will we know** it works? The **Bullseye** method (from the book *Traction*, Weinberg and Mares, 2014) answers the channel question: brainstorm all channels in the outer ring, test three cheaply in the middle ring, and put almost all effort into the one winner at the centre.

Segments: who books, and why

segment	need	pre-order help
Friday friends groups of 4–6	a table, same day	a little
Date night couples	a good table	no
Sunday family 6–10 people	no 40-min wait	some
Office lunch team weekdays, 1 hour	eat within the hour	a lot

Target first: the green row

Figure 2.19 · Segment by situation, then pick. The office lunch team is where pre-ordering matters most, so the launch starts there. *Illustrative segments.*




Figure 2.20 · Bullseye: many channels, three tests, one winner. The winner is the channel with the lowest cost per honoured booking, here (illustratively) office admin teams. *Method: Weinberg and Mares, Traction (2014).*

4 · When it misleads. Segments drawn by age or city (“urban millennials”) look neat but don’t tell you what to build; segment by situation and need instead. A channel test that is too small gives a false “no”. And a channel that works for the first 500 customers can dry up at 50,000: the tiffin WhatsApp group worked in three PGs, and then a rival joined all thirty.

5 · In the interview. “Let me pick one segment to win first and say why, give the position in one line, and then test three channels cheaply before putting money behind one.”

CASE 2.6 · RISING CANCELLATIONS AT A FOOD APP

Root cause

Medium

IIM B Sigma PM Casebook 2026–27, pp. 86–87

Try it first. You paid for a biryani two minutes ago. What could you see on the screen that would make you cancel before the restaurant has even started cooking? List three things.

THE QUESTION

We’ve seen 4–5% increase in customer initiated cancellations in the last few weeks. How would you approach this problem?

A Before diving in, I want to understand the problem clearly.

Is the increase seen across the entire customer base, or is it concentrated in specific cities or customer segments?

Q It’s mainly in major metros and among frequent users.

A Got it. And at what stage are customers cancelling? Before the restaurant starts prepping, or later in the process?

Q Mostly before preparation begins.

A Understood. Are we aware whether this spike is recent or if cancellations have been trending upward over time? Or have we shipped any new feature recently?

- Q** Spike is recent in last 2-3 weeks, and no major product changes
- A** Good to know, one last thing, is the increase uniform across all cuisines or do cuisines which require more time such as biryani have seen an increase?
- Q** That's correct, it is more for busy cuisines like biryani and North Indian, now whom do you think will be most affected by this?
- A** Primarily working professionals who need predictable delivery during office hours. They tend to cancel quickly if the delivery swings by even 10 to 15 minutes. Students form the second largest group, but their cancellations are more price sensitive than reliability sensitive. Restaurants and riders face downstream impact, but the primary friction originates with time sensitive users.
- Q** Great. What could be driving these cancellations?
- A** Broadly, cancellations can come from four buckets:
1. Restaurant-side issues : Incorrect prep-time estimates, Kitchens overloaded, Sudden item stockouts
 2. Delivery-side issues : Slow rider assignment, Long rider travel time to the restaurant
 3. Platform or UI issues : Delivery time shown upfront is not trustworthy, Users don't know restaurant reliability
 4. External factors : Weather, Peak-hour surges, Festival traffic
- From my experience and intuition, the biggest contributors are inaccurate preparation time estimates and item unavailability after the order.
- Q** Okay, assuming that's the issue, How would you solve this?
- A** I would introduce a reliability layer for the customers on top of the existing ordering flow, to solve this
- Q** Can you please explain what all will this include?
- A** Sure, the reliability layer will have three parts:
1. Real time Kitchen Load Indicator : a live signal showing actual kitchen load based on current backlog, historic preparation times and surge or peak intensity. By letting the customer know about the delay, will set right user expectation and prevents "I didn't expect it to take this long" cancellations.
 2. Restaurant Reliability Score : a simple badge showing how consistent a restaurant meets its promised preparation time, this will help solving transparency and choosing more stable kitchens.
 3. Dynamic Item available predictor : this is kind of a predictive system that marks items as low stock, only 2 left, or likely unavailable soon.
- Q** Sounds good, what kind of risks do you see in this solution, and how will you solve them?
- A** A few key risks that I can think of are:
- Restaurants may overstate inventory
 - Low reliability scores may frustrate some restaurants
 - Real time tracking requires strong restaurant compliance
- Q** Lets say as a PM you ship this feature, how will you measure success?
- A** I would like to roll out the project in phases and by using A/B experimentation to isolate and attribute the change in cancellation to this feature and to look out for any gross negatives, the north star metric for the product will be % reduction in customer initiated cancellations, other metrics will be accuracy of preparation time indicators, NPS improvement for reliability indicator and repeat order rate.
- If cancellations drop and user trust improves, we will know the feature is working.
- Q** Great, we can close the case here, it was wonderful interacting with you

WHAT WORKED · WHAT WAS MISSING

- ✓ **Five clarifying questions, each narrowing the search:** where (metros, frequent users), when in the order (before cooking), since when (2-3 weeks, no release) and which cuisines (biryani, North Indian).
- ✓ **Named who is hurt** and why: professionals who need dinner on time.
- ✓ **Grouped causes into buckets** before guessing.
- ✓ **Planned an A/B test** to prove the fix caused the drop.
- ✗ **Ignored the best clue.** Customers cancel before cooking starts, so the cause is in the first minutes after paying (see Check).
- ✗ **The buckets overlap.** "Peak-hour surges" is external, "kitchens overloaded" is restaurant-side: the same event in two places.
- ✗ **"From my experience and intuition"** replaced a data check, and three features followed.
- ✗ **No calendar check.** A 2-3 week spike in biryani orders in metros could match a festival season such as Ramzan evenings, or a rival's offer.
- ✗ **No guardrail.** Showing a busy kitchen may cut orders along with cancellations.

Scoping kit check (RCA add-ons): Exact definition ✓ **given** · Size vs baseline: 4-5% of what? ✗ · Since when ✓ · Where it shows ✓✓ · Recent changes ✓ · Can we trust the data? ✗. "A 4-5% increase" could mean 3.0% → 3.1% (noise) or 3% → 7.5% (a crisis). Always ask.

CHECK · “THE BIGGEST CONTRIBUTORS ARE WRONG PREP-TIME ESTIMATES AND ITEMS UNAVAILABLE AFTER ORDERING” DOESN'T HOLD UP

The interviewer said customers cancel **before preparation begins**. Wrong prep-time estimates show up later, as food running late. An item that turns out to be unavailable is usually the restaurant cancelling, not the customer. Early customer cancellations point to what the customer saw in the first minutes: a long “waiting for the restaurant to accept”, an arrival time that jumped after payment, or second thoughts about the bill.

STRONGER ANSWER



Figure 2.21 · A root-cause answer in six moves. Move ④ is the one the original skipped: using the interviewer's own answers to cut the tree down. *Author's summary.*

A ① CLARIFY

Customer cancellations as a share of orders: from what level to what? And how long after payment do people cancel? One more: do we trust the tagging, or could restaurant cancellations be logged as customer ones?

Q From about 3% to 7.5% of metro orders. Most cancel within three minutes of paying. The tagging is fine.

A ② SIZE

Illustrative: 4.5 extra cancellations per 100 orders, on 20 lakh metro dinner orders a day, is 90,000 cancelled dinners a day. Each is a refund, a restaurant annoyed, and a customer who may try another app tomorrow.

A ③ TREE

I'll split by the stage of the order, which has no overlaps: before the restaurant accepts, accepted but not cooking, cooking, on the way (Figure 2.22). Your data puts us in the first two. Within them, the customer either waited too long for acceptance, saw the arrival time jump, got a price surprise, or changed their mind.

A ④ CLUES

Biryani and North Indian, metros, frequent users, 2–3 weeks, no release: that looks like a demand spike at heavy kitchens. I'd check two numbers for those restaurants: time to accept an order, and arrival time after payment against the time shown on the menu. Suppose acceptance went from 1 to 4 minutes and the arrival time jumps by 20 minutes on 40% of orders.

Why this move: every clue the interviewer gave removes branches. Using them out loud is what they are testing.

A ⑤ FIX

Show the truth before payment, not after: calculate the arrival time from the kitchen's live queue before checkout. Pause a restaurant for new orders when its acceptance time passes 3 minutes. Offer “similar biryani, 30 minutes” from a less busy kitchen.

A ⑥ MEASURE

Goal (North Star): customer cancellations within 5 minutes of payment, back to 3%. **Safety check (guardrail):** orders placed, since honest times may put some people off, and order value for paused restaurants. **Trade-off:** a few lost orders against many fewer broken promises; **net orders delivered** decides.

FRAMEWORK · ISSUE TREES AND MECE

1 · The situation. The candidate had good clues and four buckets, but the buckets overlapped and nothing was ruled out. An issue tree turns “why?” into a list of branches you can check one by one.

2 · An everyday example. Your keys are missing. Split the search: inside the flat or outside it. Inside: bedroom, kitchen, living room, bathroom. Every place is in exactly one branch, and together they cover the whole flat. Now compare a bad split: “in my bag” and “in the bedroom”. The bag may be in the bedroom, so you search it twice, and nothing covers the car.

Tree level	Missing keys	Cancellations
The question	Where are my keys?	Why did early customer cancellations rise?
First split	Inside / outside the flat	Stage: before accept / accepted / cooking / on the way
Second split	Room by room	Waited too long / time jumped / price surprise / changed mind
Prune with a clue	“I had them when I came in”: skip outside	“Before cooking”: skip the last two stages
Check a leaf	Look behind the sofa cushions	Acceptance time at biryani kitchens

3 · The framework. An **issue tree** breaks one question into branches, and each branch into smaller ones, until every leaf can be checked with data. The branches should be **MECE**: mutually exclusive (no overlaps) and collectively exhaustive (no gaps). Four reliable ways to split: by **stage** of a process, by a **formula** (orders = users × frequency), by **segment** (city, platform, new or old users), and **inside vs outside** the company.



Figure 2.22 · The cancellation tree, cut down by the interviewer's clues. Splitting by stage is MECE: every cancellation happens at exactly one stage. Two branches are ruled out by one fact, and the two red leaves share a cause. *Illustrative data for the case.*

4 · When it misleads. A complete MECE tree can have fifty leaves; the skill is pruning with data, not drawing. Category splits such as restaurant / delivery / platform / external tend to overlap because causes interact: rain (external) slows riders (delivery), which makes kitchens wait (restaurant). And the neat split of the flat into inside and outside still misses keys left in the door lock, on the boundary.

5 · In the interview. “I’ll split this by the stage of the order, because every cancellation happens at exactly one stage. Your point that people cancel before cooking lets me drop half the tree straight away.”

CASE 2.7 · ZEPTO'S DELIVERY-TIME BREACHES IN BENGALURU Root cause Hard IIM B Sigma PM Casebook 2026–27, pp. 82–83

Try it first. A 12-minute delivery now takes 15, but only in some parts of the city, and it happened suddenly a month ago. Write down the steps of one delivery and guess which step grew.

THE QUESTION

Zepto in Bangalore is observing breaches. Can you please diagnose the issue?

- A** Sure. Before I begin, I would like to ask some clarifying questions. What do we mean by breaches here? How do we measure these breaches at Zepto?
- Q** For Zepto, we measure average delivery time per order. We are seeing a rise in the same.
- A** Okay. How much is the rise observed in average delivery time per order?
- Q** We are seeing about 25% increase in the delivery times.
- A** What categories are we observing these breaches in? Is it in a particular category like grocery, stationery, etc. or across all categories?
- Q** The rise is seen across all categories.
- A** Is there a specific region of Bangalore in which we are seeing this increase or is it across Bangalore?
- Q** Good question. We are observing this in few parts of Bangalore. Other regions seem to be doing fine.
- A** Alright. Lastly, for how long have we seen these breaches? Are these gradual or sudden?
- Q** We have seen this increase for a month now. Apparently, our dashboards show a step increase all of a sudden.
- A** Okay. To begin with, I would first like to analyse some broad external factors that might impact our delivery commitments. I would then go ahead and understand the user journey of our delivery system to look for any issues at any of the steps. Does this seem okay to you?
- Q** Yes. The structure looks good. Please go ahead.
- A** Thinking of seasonality as an external factor, we might see this rise because of rains or bad weather conditions in particularly our affected areas of Bangalore. Do we have any such observation?
- Q** No. The weather is good and there has been no drastic change compared to last month across Bangalore.
- A** Okay. There are a few other factors that come to my mind. There can be some government interference affecting regional small stores that we serve from, impacting our supply chain. This rise could also be due to closing of some stores in those regions making our delivery guys take longer to procure items. Does any of this seem to be a probable reason?
- Q** These are interesting factors. But there are no such issues that have been observed and everything looks good.
- A** Alright. I would then go ahead and analyse the user journey of our delivery system to get to the root cause of the issue.

The candidate breaks the delivery journey into four steps with metrics under each; the original diagram is redrawn as Figure 2.23, labelled "redrawn from the original".

- A** I have broken down the entire journey into 4 steps and listed a few metrics that I would check to see which part might be causing delays in delivery. Before we check the metrics, I would like to clarify the context of an order pickup. Delivery people might either go first to a dark store to pick up the order and then straight to the customer location or pick up orders for multiple customer by optimizing their route and then proceed towards end delivery. Should I consider multiple orders or a single order delivery for this problem?
- Q** Great question. For simplicity, let us assume the use case of a single order delivery.
- A** Okay. So, I will keep the journey simplistic with respect to order pickup and analyse the metrics. Since the issue is observed in only a few parts of Bangalore, server-side error rates or any technical issues should not be the problem. Has any of the listed metrics across the journey been impacted for that particular region over the last month?
- Q** All the metrics look fine on the dashboard except the last one which you talked about. The average time from pick-up till delivery per order has gone up.
- A** Okay. With this understanding, I would like to deep dive into just this part of the journey to find the root cause. Also, I would like to get an idea of the kind of increase in delivery time we are observing. From what I know, Zepto gives a 10 to 15 minute delivery assurance. So, a breach of 25% means the deliveries are delayed by 3-4 minutes on an average. Does my approach seem okay?
- Q** It's great that you brought out some numbers in the discussion. The approach is okay. Please go ahead.
- A** Sure. Now, the time spent from pickup till delivery can be further split up into two parts – pickup till landmark of the location and landmark to exact address. For the first part, I can think of a few factors responsible. There might be some construction work in progress in these regions causing traffic to slow down. It might also be possible that drivers have started stopping enroute for some refuelling or refreshments. Does any of this seem to be the reason?
- Q** Good points. But, all these do not seem to be a problem in these regions. Think about how a typical end user location is in your campus at IIM Bangalore and then analyse.
- A** Sure. I'll analyse the second part of this delivery time keeping a typical IIM Bangalore customer in mind. From my experience here, a significant amount of time gets consumed in the main gate in verifying the delivery guy details and directing the driver inside the campus, which increases the delivery time. Should I go ahead with to analyse it further?
- Q** Yes. This looks good. Please go ahead and extend it to a typical customer in Bangalore
- A** Now, in Bangalore, a significant number of people live in mid-to-large-sized apartments with security guards at the main gate. Now, it is possible that these 3- 4 minutes get wasted due to delivery guys talking to the guard, entering the apartment, and finally reaching the flat or house of the customer. Are we seeing any delays with respect to dealing with guards and reaching the customer?
- Q** Yes, we are. We came to know that the apartments in these regions do not have a smooth entry process for delivery agents. Hence, our guys waste about 5 minutes in the formalities with the guards and get delayed reaching the customer. That was a very thorough analysis from your side. Thank you so much.




Figure 2.23 · The candidate's four-step breakdown, redrawn from the original. Four stages of a delivery, each with the numbers to check. Only one number moved, which let the candidate zoom into the last leg. Redrawn from the original, IIM B Sigma PM Casebook 2026–27, p. 82.

WHAT WORKED · WHAT WAS MISSING

- ✓ **Textbook clarifying:** what a breach is and how it is measured, how big, which categories, which areas, sudden or gradual.
- ✓ **Outside causes first, then the journey:** the right order.
- ✓ **Four steps, each with metrics,** then asked which one moved.
- ✓ **Turned 25% into minutes** (3–4 on a 10–15 minute promise).
- ✗ **A sudden step in a few areas means something changed on a date.** The first question should have been: what changed there a month ago?
- ✗ **Never asked if the number could be trusted.** If riders started tapping “delivered” at the flat door instead of the gate, time would rise with no real change.
- ✗ **No fix, no measure.** The case ends at the cause.

✓ **Split the last leg** into “to the landmark” and “landmark to the door”, and took the hint quickly. ✗ **Missed the knock-on.** A rider stuck at a gate is late for the next order too (see the framework).

Scoping kit check (RCA add-ons): Exact definition ✓ · Size ✓ · Since when, sudden or gradual ✓ · Where it shows ✓ **category and area** · Recent changes ✗ **skipped** · Can we trust the data? ✗ **skipped**. The two skips are exactly the two that fit a sudden, local step.

CHECK · “THE ISSUE IS IN ONLY A FEW AREAS, SO TECH CAN’T BE THE PROBLEM” **OVERSTATED**

A good instinct, stated too firmly. Tech can fail locally: a new map version for one area, a delivery zone drawn wrongly so homes get served from a farther store, or a rider-app bug on one phone model common in one hub. A safer line: “A regional pattern makes a server-wide failure unlikely; I’ll check local changes, including tech ones.”

STRONGER ANSWER



Figure 2.24 · A delivery-time diagnosis in six moves. Move ⑤ turns a local annoyance into a fleet problem, which is what makes the fix worth paying for. *Author’s summary.*

- A ① CLARIFY**
So: order-to-door time is up about 25%, roughly 12 to 15 minutes, only in some areas, as a sudden step a month ago. Two more questions. What changed in those areas around that date? And did anything change in how “delivered” is recorded?
- Q** No app changes, and the recording is the same. We don’t know what else changed.
- A ② WHAT CHANGED**
A sudden step in a few areas usually means a local event: a store moved, road works, or a new rule at residential gates. Rather than guess, I’ll break the 12 minutes into steps and compare affected and normal areas.
- A ③ BREAK DOWN**
One order (Figure 2.25): assign 0.5 + pick and pack 2 + hand to rider 0.5 + ride 6 + gate and lift 2 + handover 1 = 12 minutes. Same split for the affected areas.
- A ④ FIND THE STEP**
Only “gate and lift” moved: from 2 minutes to 5. That fits the interviewer’s finding that guards now hold riders at the gate. I’d confirm it by checking whether the delay sits in large gated societies and not in independent houses in the same area.
- A ⑤ KNOCK-ON**
A stuck rider is also late for the next order. Illustrative: the area gets 2 orders a minute at peak and a round trip took 13 minutes, so $2 \times 13 = 26$ riders are busy. With 3 more minutes per trip, $2 \times 16 = 32$ are needed. With the same 26, orders queue for a rider, and even homes without gates get slower.
Why this move: it shows the fix pays for itself across the whole area, not just in gated societies.
- A ⑥ FIX + MEASURE**
Fix: link with society gate apps so the resident pre-approves the rider at checkout; learn each building’s gate time and include it in the promise; add riders at peak until the gate time falls. **Goal (North Star):** delivery time in those areas back to 12 minutes. **Safety check (guardrail):** rider earnings per hour and time riders spend waiting unpaid. **Trade-off:** extra riders cost money now; broken promises cost customers.

FRAMEWORK · BREAKING A PROCESS INTO STEPS, AND LITTLE’S LAW


1 · The situation. “Delivery is slower” is too big to fix. “The gate step went from 2 to 5 minutes in large societies” can be fixed by Tuesday.

2 · An everyday example. You have reached college 15 minutes late every day this week. Time each step of your morning: alarm, bath, dressing, breakfast, lift, bus. Only one changed: the lift in your building is under repair, so you take the stairs and miss your usual bus. And the knock-on, which is Little’s law: a dhaba with 10 tables, where 20 customers arrive an hour and each stays half an hour, has $20 \times 0.5 = 10$ tables busy on average. If meals slow to 45 minutes, it needs 15 tables. With 10, a queue forms outside, even for people who eat quickly.

Idea	Your morning / the dhaba	Zepto in Bengaluru
Break into steps	Alarm → bath → dress → lift → bus	Assign → pack → hand over → ride → gate → door
Compare with normal	Last week’s mornings	Areas that are fine
The step that moved	The lift	The gate: 2 → 5 minutes
Little’s law	20 an hour $\times$ 0.5 h = 10 tables busy	2 a minute $\times$ 13 min = 26 riders busy
Knock-on	Slow meals make a queue for everyone	Slow gates make everyone’s order wait for a rider

3 · The framework. **Process decomposition** splits a total time (or a funnel) into steps that happen one after another, measures each step, and compares with a normal baseline. The step that changed holds the cause. **Little’s law** says that, on

average and in a steady state, **items in the system = arrival rate × time each spends**. It links a delay in one step to the capacity you need.



Little's law for riders: riders busy = orders per minute × minutes per round trip

Before

2 a minute × 13 min
= 26 riders busy

Gate adds 3 min

2 a minute × 16 min
= 32 riders needed

Still only 26 riders

6 short at peak: orders wait,
even to homes with no gate

Figure 2.25 · Find the step, then follow its knock-on. Top: the same order in a normal and an affected area; only the gate step grew. Bottom: because riders are tied up longer, the area needs 6 more riders at peak, or every order slows. *Illustrative timings consistent with the case (25% rise, ~5 minutes lost at gates).*

4 · When it misleads. Steps that happen at the same time can't be added: making tea while you shower doesn't add to your morning, and a rider riding to the store while the order is packed doesn't add either. Averages hide spread: a few 30-minute gates can hide inside a "5-minute" average. And Little's law holds for averages in a steady state, not in the first ten minutes of a sudden downpour.

5 · In the interview. "Let me break the 12 minutes into steps and compare with a normal area. Whichever step grew, I'll then check its knock-on on rider capacity using Little's law."


KEY IDEA · WHAT THESE SEVEN OPENINGS TEACH ABOUT SCOPING

Pattern	Where it showed up	The question to ask
A sudden step in one place	Zepto's breaches; cancellations in metros	"What changed there, on that date?" Then: "Can we trust the number?"
"X% increase" with no base	Cancellations up "4–5%"	"From what level to what level?"
A goal question left unanswered	Table reservations; the IPO metrics	Push once, then state your own: "I'll assume the goal is honoured bookings."
One word hiding two users	"Vendors": national brands and local bakeries	"How is this group split, and which part matters most?"
An asset the company already owns	Swiggy's Dineout, ignored in a reservations launch	"What do we already have here?"

Asking the right one of these in the first two minutes changed the whole answer in every case above. None needs more than one sentence.

Zomato and Blinkit: profit from food, growth from ten minutes

Walkthrough (illustrative persona). 8:40 pm. Kavya opens Zomato. The top row says “Order again” and shows the biryani she had two Fridays ago. She taps it, sees “arrives in 28–33 min”, and the bill: ₹450 food, ₹25 delivery, ₹15 platform fee. She pays by UPI. A map shows the restaurant preparing, then a rider’s dot moving towards her. At the door she reads out a four-digit PIN. The whole visible journey takes four taps; Figure 2.26 shows the relay behind them.




₹505 comes in (₹490 from Kavya, ₹15 of ads); ₹445 goes out. Zomato runs the whole relay but cooks and rides nothing.

Figure 2.26 · What happens behind Kavya’s four taps. Six parties, thirteen steps. Zomato holds the money for a few days and makes every match, but the food and the ride belong to others. Solid arrows are requests, dotted arrows replies, green arrows money. *Simplified from Zomato’s public descriptions of ordering, dispatch and payouts; amounts illustrative.*

Key features, and why they’re built that way

Feature	User problem	Why it’s designed this way	Metric it moves
“Order again” and dish search	Choosing dinner takes too long	Most orders are repeats, so the fastest path is the last good meal; search works on dishes, not just restaurants	Time from open to payment; conversion
Arrival time as a range, live map	“Where is my food?”	A range is kept more often than a single number; the map cuts “where is my order” calls to support	On-time %; support contacts per 100 orders
Zomato Gold	Fees add up for frequent users	A paid membership with free delivery above a basket size; members order more often	Orders per member vs non-member
Platform fee	Zomato: thin margin per order	A flat fee in rupees, not a %, shown before payment	Revenue per order (guardrail: conversion)
Restaurant ads	Restaurants: “customers can’t find us”	Pay for a higher place in search and on the home screen	Ad revenue as % of order value
Blinkit’s range beyond groceries	“I need a charger / a gift / printouts now”	Electronics, toys and beauty raise the basket and give a reason to open the app daily	AOV; orders per customer

The second app. 10:52 pm, Gurugram. Rohan, a student, realises there is no milk for the morning and his laptop charger has died. On Blinkit he adds both, pays ₹1,340, and a rider is at his door at 11:03. The store 1.8 km away never saw him; its picker saw only a list of shelf locations (Figure 2.27).



Since Blinkit owns the stock (1P), step ⑩ is its money at risk: unsold milk is its loss.

Figure 2.27 · Behind an 11-minute Blinkit order. Steps ①–⑥ are the customer's 11 minutes; ⑦–⑩ happen overnight and decide whether tomorrow's order will be in stock. Simplified from Eternal shareholder letters (2025–26); amounts illustrative.




Figure 2.28 · The core loop. More orders in a neighbourhood make riders busier and delivery cheaper, which brings more orders. The same loop runs in Blinkit's stores. Author's model.

Order value, Apr–Jun 2026 (₹ Cr) · adjusted EBITDA

Blinkit	17,132
	+₹102 Cr, first profit
Food	10,769
	+₹606 Cr (5.6%)
District	~3,200
	–₹65 Cr

Figure 2.29 · Eternal's three consumer businesses. Blinkit is now the largest by order value; food still makes most of the profit. Eternal Q1 FY27 shareholder letter (Jul 2026); District derived from total B2C NOV of ₹31,120 Cr.

Business model in one line. Food earns a commission, customer fees and ads on each order, and turns about 5–6% of order value into profit. Blinkit now owns much of its stock, so its revenue counts the full price of goods (which is why Eternal's revenue jumped 182% to ₹20,211 Cr), and it earns the product margin plus brand ads. Hyperpure sells supplies to restaurants.

NORTH STAR Orders delivered per week (food) and order value per store per day (Blinkit).	INPUTS Monthly transacting customers · orders per customer · top restaurants online at 8 pm · store fill rate · share of mature stores.	GUARDRAILS Cancellations · on-time % · rider earnings per hour · inventory written off as % of sales.
--	---	---

Moat, in plain words. Density. Zomato has more restaurants and more orders in most neighbourhoods, so its riders are busier and its delivery cheaper. Blinkit has the most dark stores (2,443 in Jun 2026), and a mature store handles about ₹8.3 lakh of orders a day, which a newcomer's half-empty store can't match. Both cracks are the same: restaurants and brands list on every app, and customers switch for a ₹50 coupon.

2025–26 moves. Renamed the parent Eternal (2025). Blinkit moved to owning stock after Eternal became Indian-owned and controlled. District, a separate app for dining out, films and events. Bistro, a 10-minute café app. Zomato's own 10-minute food experiments shut. Blinkit posted its first adjusted profit (Apr–Jun 2026).

STRENGTHS

- ✓ Leader in food and in quick commerce, both profitable at the operating level
- ✓ Dense store network and a store economics playbook rivals copy
- ✓ Food profits fund experiments without raising money

WEAKNESSES

- ✗ Food growth has slowed to about 20% a year
- ✗ Owning stock adds waste and working capital risk
- ✗ Restaurant goodwill: the CCI report and commission complaints


What I'd build next

- 1. Families order dinner separately on four phones → a shared family cart with one payer and one delivery → AOV and orders per household.
- 2. On rain nights promised times break and cancellations spike → a “rain mode” that widens the time range before payment and hides restaurants beyond 5 km → cancellations on rain nights vs dry nights.
- 3. Blinkit stores waste fresh stock that is close to expiry → a 7–10 pm “use today” shelf at a discount, shown only near the store → inventory written off as % of sales.

Swiggy and Instamart: one member, one fleet, many services

Walkthrough (illustrative persona). Neha, 34, a consultant in Pune, pays for Swiggy One. At 8 pm she orders dal makhani and rotis for ₹380; the membership makes delivery free. While waiting she remembers breakfast, taps the Instamart tab in the same app, and adds eggs and bread for ₹240. Dinner arrives at 8:34, groceries at 8:47. It feels like one shop. Behind it, two businesses share one customer, one wallet and often one rider (Figure 2.30).


Feature	User problem	Why it's designed this way	Metric it moves
Swiggy One	"I pay delivery fees on every app"	One membership across food, Instamart and Dineout makes Swiggy the default for everything	Services used per member; orders per member
Bolt (since Oct 2024)	"I want food in 10–15 minutes"	Only restaurants within about 2 km and dishes that are quick to make, so speed comes from distance, not rushing	Share of orders under 15 min; repeat rate
Instamart megapods	"They never have what I need"	Larger dark stores with 10,000+ items for bigger, planned baskets	Fill rate; AOV (about ₹700)
Basket-size offers	Swiggy: small baskets lose money	Discounts unlock only above a basket value	AOV; contribution per order
Dineout	Paying the bill and finding a deal when eating out	Book a table, pay the bill in the app, get a restaurant-funded discount	Dine-out bills paid through the app



One rider, one member, two businesses: the shared fleet and the One membership are what tie food to Instamart.

Figure 2.30 • Behind Neha's evening. Food and Instamart orders use different suppliers but the same fleet and the same membership. The rider in step 9 is the one who has just delivered the dinner. Simplified from Swiggy's public descriptions of its shared delivery fleet and Swiggy One; amounts illustrative.

Order value, Apr–Jun 2026 (₹ Cr), growth, margin



Group net loss: ₹791 Cr in the quarter

Figure 2.31 · Two businesses, two stages. Food pays the bills; Instamart has just stopped losing money on each order but still loses money after fixed costs. *Swiggy Q1 FY27 results (Jul 2026) as reported by Business Standard and Entracker; Q4 FY25 margin from the company.*




Figure 2.32 · The growth loop. Swiggy's bet is that people who use two services stay longer and cost less to serve. *Author's model of Swiggy's stated strategy.*

NORTH STAR	INPUTS	GUARDRAILS
Monthly transacting users who use two or more services.	Food orders per user · Instamart AOV · fill rate · One members · rider orders per hour.	Contribution per order (both) · on-time % · membership cancellations.

Moat, in plain words. A shared fleet and a shared membership: every service Swiggy adds makes both cheaper per order. It is weaker than it sounds, because Zomato and Blinkit together match it in most big cities and Blinkit has more stores.

2025–26 moves. Instamart chose profit over size: it dropped millions of unprofitable, small-basket users and pushed bigger baskets, reaching contribution breakeven in May 2026. Bolt grew as a share of food orders. Snacc, a separate 10-minute snacks app, was tried in 2025. Advertising spend stayed high (about ₹1,160 Cr in Apr–Jun 2026) to hold share.

STRENGTHS

- ✓ Highest basket size in quick commerce
- ✓ One app, one fleet, one membership
- ✓ Food business profitable and growing

WEAKNESSES

- ✗ Number two in food, number three in quick commerce
- ✗ Still loss-making as a group
- ✗ Fewer dark stores than Blinkit


What I'd build next

1. Members don't notice Instamart while waiting for food → an "add groceries, same rider" card on the order-tracking screen for pods under 1 km away → share of food orders with a grocery add-on.
2. Bolt users can't tell what is fresh-cooked → show "made to order" vs "ready-made" on each Bolt dish → Bolt repeat rate and ratings.

Zepto: speed first, and then the bill for it

Walkthrough (illustrative persona). Ira, 22, is studying late in a Mumbai PG. She opens Zepto for a snack. The top row shows a new flavour of chips marked "Ad". She adds it, then a cold coffee from "Zepto Café", and pays ₹310. Nine minutes later one bag arrives with both. The chips maker paid Zepto before Ira saw the chips; the coffee was brewed at a counter inside the same dark store (Figure 2.33).

Feature	User problem	Why it's designed this way	Metric it moves
10-minute promise	"I forgot something"	Dense stores close to homes and picking in about 2 minutes	Delivery time; orders per customer
Zepto Café	A snack or coffee with the groceries	Small kitchens inside dark stores add a high-margin item to a small basket	Café attach rate; margin per basket
Ads for brands	Brands: "we need to be seen where shoppers decide"	Sponsored rows and banners inside search and categories	Ad revenue as % of order value
Membership pass	Small baskets attract fees	A monthly pass removes delivery and handling fees	Orders per member



The café adds a high-margin item to a small basket; the ad pays Zepto before anything is sold.

Figure 2.33 · Behind Ira's snack order. Two things raise the value of a small basket: a café item with a high margin (③–⑤) and an ad paid for by a brand (⑧). Simplified from Zepto's public descriptions of Zepto Café and its ads platform; amounts illustrative.




Figure 2.34 · A dark store's life. Contribution margin climbs as orders per day rise. Opening hundreds of stores at once makes the whole company look worse even as older stores improve. Illustrative shape, consistent with company statements that new stores take 6–12 months to break even.

KEY IDEA · WHY FAST GROWTH LOOKS LIKE BIG LOSSES

Imagine stores that each lose ₹3 lakh a month in their first six months and make ₹1 lakh a month after that. A company that opens 50 new stores this quarter reports a bigger loss than last quarter, even though every old store is improving. So ask two questions of any quick-commerce loss: **what share of stores are under six months old**, and **what do the mature stores earn**? Blinkit's first profit came when its expansion slowed and most stores had matured.

Business model and 2026 reality. Zepto earns the margin on goods sold through its stores, customer fees, café margins and brand ads. Press reports put FY26 revenue at about ₹22,600 Cr and the loss at about ₹5,900 Cr (not yet checked against filings). It moved its holding company from Singapore to India in 2025, filed for an IPO, and in July 2026 put the IPO on hold and raised money privately instead.

NORTH STAR Weekly orders from repeat customers.	INPUTS Orders per store per day · share of stores over 8 months old · AOV · fill rate.	GUARDRAILS Contribution per order · cash burn per month · stock written off.
---	--	--

Moat, in plain words. Thin. Speed is copyable by anyone with stores in the same streets, and Blinkit, Instamart, Flipkart Minutes and Amazon Now all have them. What Zepto had was execution speed and a young brand; what it lacks is a profitable second business to pay for the fight.

STRENGTHS

- ✓ Fastest-growing challenger; strong with young urban users
- ✓ Café and ads raise margin per basket

WEAKNESSES

- ✗ Large losses; IPO on hold
- ✗ No profitable core business, unlike Eternal or Swiggy

What I'd build next


1. Young stores lose money until orders build → open new stores only where a nearby store already fills 1,000 orders a day, and move its far orders to the new store → months to breakeven per new store.
2. Brands can't see which stores sell their products → self-serve store-level sales and ad results for brands → ad revenue per brand; brands advertising again next quarter.

BigBasket: the planned weekly shop in a ten-minute world

Walkthrough (illustrative persona). Mrs Iyer, 58, in Chennai, orders the week's groceries every Sunday night: vegetables, rice, dal, oil and cleaning supplies, ₹2,150 in all, for the 7–9 am slot. She trusts BigBasket's vegetables

more than the corner shop's and earns NeuCoins she spends on Tata products. On Wednesday she runs out of curd and uses BB Now, which arrives in 18 minutes. One app, two machines (Figure 2.35).

Feature	User problem	Why it's designed this way	Metric it moves
Scheduled slots	A big weekly shop with fresh produce	Orders known a night ahead: buy produce to demand, pick overnight, one van route for 20 homes	Cost per order; freshness complaints
BB Now	"I ran out of curd"	10–20 minute delivery from dark stores, inside the same app	Share of households using both modes
Own brands (Fresho, BB Royal)	Trust and price	Graded produce and staples under BigBasket's own label, at a better margin	Own-brand share of sales
Tata Neu and NeuCoins	Loyalty that is worth something	Coins earned here can be spent at Croma, Taj or Westside	Customers active on two Tata brands



Scheduled slots let BigBasket buy fresh produce against known demand and fill a van with 20 orders on one route.

Figure 2.35 · Behind a scheduled weekly basket. Knowing tomorrow's orders tonight lets BigBasket buy produce to demand (2–3), pick overnight, and put 20 homes on one van route, which is cheaper per order than any 10-minute ride. Simplified from BigBasket's public descriptions of its supply chain; amounts illustrative.




Figure 2.36 · Paying to become fast. Revenue barely moved while losses widened: the cost of building quick commerce on top of a scheduled business. FY26 filings as reported by MediaNama and Inc42 (Sep 2026); FY25 values derived from reported growth rates.

REAL-LIFE EXAMPLE · WHY A VAN BEATS A SCOOTER, SOMETIMES

Illustrative: a van costs ₹2,400 for a morning route of 20 homes, so ₹120 an order on baskets of ₹2,000+ (6% of the basket). A 10-minute rider costs ₹45 on a ₹500 basket (9%). Per rupee of groceries the van is cheaper; per order the scooter is. That is why BigBasket keeps both: slots for the weekly shop, BB Now for top-ups. The risk is that the top-ups slowly replace the weekly shop, and the cheaper machine goes idle.

NORTH STAR	INPUTS	GUARDRAILS
Monthly order value from households that buy weekly.	Weekly buyers · BB Now attach · own-brand share · slot on-time %.	Fresh-produce complaints · loss per order in BB Now stores.

Moat, in plain words. Fresh-produce sourcing and grading, built over a decade, and the trust of planned, older, family shoppers. The weakness: it came late to speed, grows slowly (6%), and its app is opened weekly while rivals' are opened daily.

STRENGTHS

- ✓ Best in fresh produce and big planned baskets
- ✓ Tata's balance sheet and loyalty network

WEAKNESSES

- ✗ Slow growth, widening loss
- ✗ Seen as slower and less exciting by young users

What I'd build next

1. Weekly shoppers forget items and turn to rivals mid-week → a "top up before Wednesday" reminder built from each household's usual run-out day, delivered by BB Now → share of weekly buyers with a mid-week BB Now order.

2. Produce quality is BigBasket’s edge but invisible in a 10-minute app → a “picked today, graded A” tag with the farm region → fresh-category conversion and complaints.

magicpin and ONDC: food delivery on an open network

Walkthrough (illustrative persona). Sameer, 27, in Delhi, noticed that his usual paneer roll costs less on magicpin than on the big apps: same shop, same roll. He searches, orders for ₹240, and a rider from a separate delivery company brings it. What he doesn’t see is that four different companies handled his order, all speaking one open protocol run by ONDC, the government-backed Open Network for Digital Commerce (Figure 2.38).

What the restaurant keeps from a ₹400 order (illustrative)

Big food app



ONDC via magicpin




red: commission, fees and ads the restaurant pays

Figure 2.37 · Why restaurants join. Lower fees let them keep more, or price the same roll lower. Delivery is often paid separately by the customer. *Illustrative; commissions of 18–25% on large apps and single-digit fees on ONDC are commonly reported (2024–26).*

CHECK · “ONDC WILL BREAK THE FOOD-DELIVERY DUOPOLY” OVERSTATED

At about 1.5 lakh orders a day in late 2024, magicpin’s ONDC food business was a few per cent of the big two’s volume. Lower fees help restaurants, but customers choose on range, reliability and habit, and those still sit with Zomato and Swiggy. A better claim: ONDC puts a ceiling on how far commissions can rise, because restaurants now have somewhere else to go.

Feature	User problem	Why it’s designed this way	Metric it moves
Lower prices	“Same food, why pay more?”	Lower network fees are passed on as lower menu prices	Share of orders from price-sensitive users
magicNOW	“I want it fast”	15-minute delivery from kitchens very close to the user	Share of orders under 15 minutes
magicPoints	A reason to come back	Loyalty points carried over from its discovery days	Repeat rate
Its own delivery fleet	Riders from other companies are less reliable	A fleet it controls for the busiest areas	On-time %



No single company owns the customer, the restaurant and the rider: each role is a separate app speaking one open protocol.

Figure 2.38 · Behind an ONDC food order. The buyer app (magicpin), the restaurant’s seller app and the delivery company are separate businesses. ONDC sets the rules and passes messages; it holds no money and no customers. *Simplified from ONDC’s published network design (Beckn protocol); amounts illustrative.*

Business model and scale. magicpin earns fees as a buyer app, delivery fees and brand deals. It became the largest food buyer app on ONDC, at about 1.5 lakh orders a day (Oct 2024). ONDC’s own incentives, which once paid part of each order, have been cut back, and that tests how many customers stay without them.

NORTH STAR Weekly repeat orders without any subsidy.	INPUTS Restaurants live at 8 pm · price gap vs big apps · delivery time.	GUARDRAILS Order failures across apps · refund time · complaints.
--	--	---

Moat, in plain words. Very little of its own: any buyer app can plug into the same restaurants on ONDC. Its edge is being early, a loyalty base and a delivery fleet. The deeper question is whether an open network can match the reliability of one company that controls the whole chain: when an ONDC order fails, whose fault is it, and who refunds?

STRENGTHS

- ✓ Real price advantage for the same food
- ✓ Restaurants actively want an alternative

WEAKNESSES

- ✗ Smaller range and less reliable delivery than the big two
- ✗ Growth leaned on ONDC incentives

What I'd build next

1. Users can't tell who is responsible when an order goes wrong → a single magicpin guarantee: refund in 10 minutes, whichever company failed → repeat rate after a failed order.
2. The price gap is the reason to switch, but users can't see it → show “₹38 less than on other apps” on each dish, from restaurants' own price data → conversion on dishes showing the gap.

Side by side

	Zomato + Blinkit	Swiggy + Instamart	Zepto	BigBasket	magicpin (ONDC)
Core promise	Most restaurants; 10-min everything	One membership for all	Fastest snacks and top-ups	Fresh, planned weekly shop	Same food, lower price
Who owns the stock	Restaurants; Blinkit mostly itself	Restaurants; sellers in pods	Sellers and itself	Itself (Tata-owned)	Restaurants
Scale (Apr–Jun 2026)	Food ₹10,769 Cr; Blinkit ₹17,132 Cr	Food ₹9,490 Cr; Instamart ₹7,907 Cr	Private; FY26 loss ~₹5,900 Cr	FY26 revenue ₹10,498 Cr	~1.5 lakh orders a day (2024)
Profit	Both profitable (adjusted)	Food profitable; Instamart at breakeven	Loss-making	Loss-making	Depends on incentives
Moat	Density, store count	Shared fleet, membership	Speed (copyable)	Produce sourcing	Lower fees (copyable)
Biggest risk	Waste from owning stock	Squeezed between two leaders	Cash	Too slow, too late	Reliability





Figure 2.39 · Speed against range. Scheduled shopping owns the top-left, classic quick commerce the bottom-right. The race of 2025–26 is to the top-right: stores big enough for a weekly shop and close enough for ten minutes. Positions are the author's judgement from company disclosures and store formats, 2026; not to scale.

UX design: from “users are leaving” to a screen that works

A food app’s data shows that 18 of every 100 people who reach the bill screen leave without paying. The design team wants a new checkout; engineering wants to know what to build; the business wants the fee to stay. **UX (user experience) design** is the discipline of finding out why people struggle with a product and designing it so they don’t. For a PM it is less about drawing screens and more about a process that stops the team from building the wrong thing.

Design thinking: five stages, used as a loop

Design thinking is a problem-solving process that starts from the user’s experience instead of from a feature idea. The common version (Stanford d.school) has five stages: **Empathise** (understand users), **Define** (state the problem sharply), **Ideate** (many ideas), **Prototype** (a cheap version) and **Test** (watch real users try it). It is a loop: a test usually sends you back to define the problem better.




Two days of prototype and five users can kill a bad idea before two months of engineering are spent on it.

Figure 2.40 · Design thinking on the checkout problem. Each stage produces something concrete: an observation, a sentence, a short list, a clickable screen, a result. The dashed arrow is the most common path in real life. *Process: Stanford d.school; example illustrative.*

Research: what to learn, from how many people

Empathise uses interviews (ask about the last real time, not opinions), observation (watch someone order dinner), surveys (how many feel this) and data (where people drop off). The notes pile up quickly, so teams group them with **affinity mapping** (the KJ method, after Jiro Kawakita): write each observation on a sticky note, cluster similar notes silently, then name each cluster. The clusters become the problems worth defining.



Eight interviews → notes → three themes

Figure 2.41 · Affinity mapping. The biggest cluster is not always the most important, but it is where to look first. *Illustrative notes.*




Figure 2.42 · Why five test users is usually enough. If each user hits a given usability problem with a 31% chance, five users reveal about 84% of such problems. Better: three rounds of five than one round of fifteen. *Nielsen and Landauer (1993); Nielsen, “Why you only need to test with 5 users” (2000).*

Personas that help, and ones that don’t. A **persona** is a short, research-based portrait of one type of user. A useless persona lists age, city and hobbies (“Riya, 26, likes travel”). A useful one describes a situation and a goal: “Riya orders dinner at her desk three nights a week, has 15 minutes, and cares most about arriving on time.” The second tells a designer what to put on the first screen.

Prototypes: spend as little as the question needs

Fidelity rises →

cost and time rise with it

Paper sketch
1 hour · “is this the idea?”

Wireframe
1 day · “is the layout
and order of steps right?”

Clickable (Figma)
2–3 days · “can users
finish the task?”

Coded beta
2–6 weeks · “does it move
the metric with real
orders and real money?”

Figure 2.43 · The prototype ladder. Pick the cheapest step that answers the question you have. Teams waste most time by jumping to a coded beta to answer a layout question. *Author’s summary of common practice.*

Nielsen’s ten heuristics, on a food-app checkout

Usability heuristics are rules of thumb for spotting problems without a user test. Jakob Nielsen’s ten (1994) are the standard list. Walk the screen with each one and note violations, then fix by **severity × how often users hit it**.

Heuristic, in plain words	Checkout: done well	Checkout: a violation
1. Show what is happening	“Restaurant is preparing · rider 1.2 km away”	A spinner for 40 seconds after paying
2. Speak the user’s language	“Arrives 8:55–9:05 pm”	“Order state: DISPATCH_PENDING”
3. Easy undo and exit	Edit the cart from the bill screen	Back button empties the cart
4. Be consistent	“Add” always in the same place and colour	“Proceed” on one screen, “Continue” on the next
5. Prevent errors	Warns when the address is 12 km from the pin	Lets you pay for a closed restaurant
6. Recognition over recall	Saved address and last payment method shown	Retype the flat number each time
7. Shortcuts for regulars	One-tap “order again”	Eight taps to repeat last Friday’s order
8. Keep it minimal	Bill shows total first, details on tap	Five banners above the Pay button
9. Help recover from errors	“Payment failed; money not taken; retry”	“Error 502”
10. Help when needed	“Why this fee?” link next to the fee	Help buried under Profile → Settings

WHEN IT BREAKS · GOOD UX FOR THE COMPANY, BAD FOR THE USER

Design can be used against users: a fee revealed only at the last step, a donation box ticked by default, a cancel button hidden in grey. These are **dark patterns**. India’s Central Consumer Protection Authority notified guidelines in November 2023 that name 13 of them, including **drip pricing** (revealing charges bit by bit) and **basket sneaking** (adding items or charges without asking), and treat them as unfair trade practice. A PM who ships one can win a quarter’s metric and lose the regulator and the customer. Other ways UX work breaks: “design thinking theatre”, with workshops and sticky notes that never change a decision; trusting what users *say* over what they *do*; and using five-user tests for questions of preference or price, which need numbers.

KEY IDEA · THE PM’S UX DECISIONS

- **How much research before building?** More when a mis-take is costly or hard to undo.
- **Which prototype?** The cheapest that answers the question.
- **Which problems first?** Severity × frequency, not whichever the loudest stakeholder saw.
- **Where is the line?** Nudges that help users are fine; hiding costs is not.

IN THE INTERVIEW

In any design case, one sentence shows UX maturity: “Before building, I’d watch five users try a clickable prototype of the new checkout, fix what they struggle with, and only then A/B test it on real traffic.” If asked about fees, add: “and every fee is shown before the bill, never added at the end.”

System design II: where an order lives, and why the map never freezes

IPL final, 8:30 pm. Orders on a food app run at three times a normal Sunday. In the same minute, lakhs of people open menus, thousands pay, and lakhs of riders' phones send their location every few seconds. The payment must never be lost; the menu must open in a blink; the rider's dot must keep moving. These three jobs need three different kinds of storage, and the choice between them is one of the most common PM-engineering conversations.




Figure 2.44 · One food app, simplified. The order service saves the order and payment in a SQL database, then drops an “order placed” event on a queue so that dispatch, notifications and invoices happen in the background. Menus come from a cache; rider locations never touch the order database. Author's simplified architecture, typical of large delivery apps; not any one company's design.

Databases: SQL for money, NoSQL for flexible things

A real situation. Kavya pays ₹490. Two things must be saved together: “payment received” and “order created”. If the app saves one and crashes before the other, she has paid for food nobody will cook. A **SQL (relational) database** stores data in tables with fixed columns (orders, customers, payments) and supports **transactions**: a group of changes that all happen or none do. The guarantee has a name, **ACID**, which in plain words means complete, consistent, not mixed up with other people's changes, and permanent once saved. A **NoSQL database** stores flexible documents (a menu with dishes, variants and add-ons that differ by restaurant) and spreads easily across many machines, usually by relaxing some of those guarantees.

	SQL (e.g. PostgreSQL, MySQL)	NoSQL (e.g. MongoDB, Cassandra, DynamoDB)
Shape of data	Tables, fixed columns, links between tables	Documents or key-value pairs, shape can vary
In a food app	Orders, payments, refunds, restaurant payouts	Menus, reviews, user activity, product catalogues
Strength	All-or-nothing changes; exact numbers	Huge volumes; easy to add new fields and machines
Where it strains	Very high write volumes on one machine	Money: harder to guarantee two changes happen together

Indexes: why one query takes 2 ms and another 2 minutes. Kavya taps “My orders”. The orders table has 200 crore rows. Without help, the database reads every row looking for hers. An **index** is like the index at the back of a text-book: a sorted list of customer IDs pointing to their rows. Finding one ID in a sorted list of 2 billion takes about 31 halving steps ($2^{31} \approx 2.1$ billion), not 2 billion reads. The catch: every index must be updated on every new order, so indexes speed up reads and slow down writes. When a PM asks for a new filter (“orders by restaurant, by week”), the real engineering cost is often a new index on a very large table.

One order, eight states

A real situation. Kavya taps “Cancel” at 8:52 pm. Does she get her ₹490 back? It depends on where the order is. Inside the database an order is a **state machine**: a fixed list of states, and a fixed list of allowed moves between them. Every move is saved with a timestamp. That log is what made the cancellations case answerable: “cancelled before preparation began” is simply a count of orders that moved to Cancelled from Placed or Accepted.





Figure 2.45 · Kavya's order as a state machine. Grey timestamps are the log that analysts query. The red dashed moves are where money and blame get decided; they are product policy written into code. *Illustrative states and refund policy, typical of food apps.*

When one database is not enough: replicas and shards

Replicas. Most requests only read data ("show my orders"), and a few write it ("place an order"). So teams keep one **primary** database that takes all writes and several **read replicas**, copies that serve reads. Copies arrive a fraction of a second late (**replication lag**). That is why an app sometimes shows "no orders yet" one second after you paid; the fix is to read your own fresh order from the primary. **Shards.** When writes outgrow one machine, the data is split across several databases by a key, each piece a **shard**. Illustrative: at peak each order writes about 10 rows (the order, items, state changes, payment), so 600 orders a second is 6,000 writes a second; four shards take about 1,500 each. The key decides everything (Figure 2.46).

Replicas: one writer, many readers



dashed: copies arrive ~0.1–1 s late, so read your own just-placed order from the primary

Shards: split by which key?

By customer ID: load spreads evenly



...but "all orders of one restaurant" must ask all four

By restaurant ID: a viral biryani place overloads one shard



hot shard: slow for every restaurant stored on it

No key is best for every question; choose by the busiest one.

Figure 2.46 · Scaling reads with replicas, and writes with shards. Replicas trade a moment of staleness for many more reads. Shards trade simple queries for more writes, and the choice of key decides which queries become slow. *Illustrative load; common patterns in large consumer apps.*

Caching: the answer you already have

A **cache** is a small, very fast memory that keeps copies of answers the system is asked for often, like the chef's counter that holds today's most-ordered sauces instead of fetching them from the store room each time. Menus are ideal: read lakhs of times, changed a few times a day. The number that matters is the **hit rate**, the share of requests the cache answers on its own.

50,000 menu views a second at peak

hit rate	reads reaching the database
80%	10,000/s
95%	2,500/s
99%	500/s

average wait = hit rate × 1 ms + misses × 20 ms

80%: 4.8 ms · 95%: 2.0 ms · 99%: 1.2 ms

80 → 95% cuts database load four times

Figure 2.47 · Cache hit-rate arithmetic. Small changes near 100% make big differences to the database behind it. *Illustrative load; typical memory and database response times.*

WHEN IT BREAKS · THE STALE MENU

The restaurant runs out of mutton biryani at 9:10 pm. The cache still shows it until its copy expires. Customers keep ordering a dish that doesn't exist, and the restaurant rejects each order. The fixes are engineering choices with product costs: a short expiry (more database load), or clearing the cached copy the moment the restaurant marks the dish sold out (more moving parts). The PM decides which data may be a few minutes old (ratings, photos) and which must be live (stock, prices, fees).

Expiry, and the stampede. Every cached copy has a **TTL** (time to live), say 5 minutes, after which it is fetched again. On a big night, if a viral restaurant's cached menu expires, the next 5,000 requests that second all miss at once and hit the database together: a **cache stampede**. The usual fixes: let one request refresh the copy while the rest are

served the old one for a second, and spread expiry times so that thousands of entries (every promo banner set to go live at 8 pm, say) don't expire in the same second. For a PM this explains a familiar incident: "menus stopped loading exactly when the 8 pm offer went live."

Queues: say "order placed" now, do the rest in the background


When Kavya pays, the app needs to do five things: save the order, find a rider, notify the restaurant, send an SMS, and create an invoice. If it did them one after another before answering, she would stare at a spinner for seconds, and one slow SMS provider could block the whole order. Instead the order service saves the order, puts an "order placed" message on a **queue**, and answers her immediately. Separate programs (**workers**) take messages off the queue and do the other jobs in their own time. This is **asynchronous processing**. A kitchen works the same way: the waiter hangs your ticket on the rail and returns to the next table; the cooks pick tickets up as they free up.

LIGHT MATHS · HOW MANY WORKERS FOR THE IPL FINAL?

Normal evening: 50 orders a second. Each order needs 0.2 seconds of background work. So $50 \times 0.2 = 10$ **workers** busy on average; run about 15 to leave room. IPL final: 150 orders a second needs 30 busy workers. With only 20, the workers can clear $20 \div 0.2 = 100$ orders a second, so the queue grows by 50 every second: **30,000 orders** waiting after ten minutes. Customers still see "order placed", but their SMS, and worse, their rider search, run late. That is why rider dispatch gets its own queue and priority, ahead of invoices and analytics.

Live tracking: 75,000 updates a second

Illustrative: 3 lakh riders online at peak, each phone sending its location every 4 seconds, is $3,00,000 \div 4 = 75,000$ **writes a second**. Writing each into the SQL orders database would swamp it. So location pings take a separate path (Figure 2.48): a queue, a fast in-memory store that keeps only each rider's latest point, and a push to the customer's app over a connection that stays open.




The order database is never touched: 75,000 pings a second go to a queue and a fast in-memory store instead.

Figure 2.48 · Behind the moving dot. One ping becomes an overwrite in memory and two pushes to the customer. History, if needed for disputes or better arrival-time models, goes to cheap storage later. *Author's simplified design, typical of ride-hailing and delivery apps.*

The PM's dial. Pinging every 2 seconds makes the dot smoother, doubles the writes to 1.5 lakh a second and drains riders' phone batteries faster. Pinging every 10 seconds saves both, and the dot jumps across the map, which makes customers call support. Most apps settle between 3 and 5 seconds while a rider is on an order, and much less often while idle.

Finding the nearest free rider in milliseconds

A real situation. Kavya's biryani will be ready at 8:58, and dispatch must pick a rider now. Measuring the distance from the restaurant to all 3 lakh riders, for each of 50 new orders a second, is 1.5 crore calculations a second, mostly about riders on the other side of the city. Instead the map is cut into small grid cells, each with a short code; a common scheme, **geohash**, gives nearby places codes that start the same way. The system keeps, in memory, a list of free riders in each cell. To find candidates it looks only in the restaurant's cell and the 8 around it (Figure 2.49).



How dispatch narrows 3 lakh riders to a handful

- the restaurant's cell
- its 8 neighbours: search only these 9
- free riders found: about 7, not 3 lakh
- riders outside the search: ignored

Then rank the 7 candidates by

- ① time by road, not straight-line distance (the blue river: the ringed rider is 600 m away, but 4 km by road)
- ② can they batch this with an order they already carry?
- ③ acceptance rate and rating ④ fairness: who has waited longest

No rider in 9 cells? Widen to the next ring, or raise the pay for this order.

Figure 2.49 · Nearby search with a grid. Looking in 9 cells instead of the whole city turns crores of distance checks into a few dozen. The ranking rules after that are product choices: speed for the customer, fair earnings for riders, cost for the company. *Illustrative; grid-based nearby search is standard in ride-hailing and delivery dispatch.*

Matching one at a time, or in batches. The simplest dispatch gives each order to the best free rider the moment it arrives. Many apps instead collect the orders of the last 20–30 seconds and match them all together, because the best rider for order A may be the only good rider for order B. Batch matching lowers total travel across the city, at the cost of a few seconds' wait before any rider is chosen. When a PM asks “why did it take 40 seconds to assign a rider?”, this is often the answer, and on a Friday at 8:30 pm it is usually the right trade-off.

WHEN IT BREAKS · WHAT GOES WRONG ON A BUSY NIGHT

- **Queue backlog:** orders show “placed” but no rider is found for minutes (the worker maths above).
- **Hot spot:** one viral restaurant gets 5,000 orders in an hour and the database machine holding its data overheats while others sit idle.
- **Out-of-sync screens:** the payment succeeded but the order screen still says “pending”, because two stores haven't caught up. Tell the user the truth (“payment received, confirming your order”) rather than a scary error.
- **Stale cache:** sold-out dishes still orderable (see above).

Exact now, or correct soon?

Keeping every copy of every number identical everywhere, instantly, is slow and expensive. So engineers ask of each piece of data whether it needs **strong consistency** (everyone sees the same, latest value, always) or can live with **eventual consistency** (copies agree after a short delay). The answer is a product decision, because only the PM knows what a stale value costs the user.

Data	How fresh?	Why, in user terms
Payment status, order state	Exact, now	A wrong “unpaid” means a double charge or no food
Dish sold out, fees, prices	Within seconds	Stale means rejected orders, or a price that changes at checkout
Rider's dot on the map	A few seconds	Smoothness matters more than exactness
Restaurant rating, review count	Minutes	Nobody notices 4.3 becoming 4.4 an hour late
Dashboards, “orders today”	Minutes to hours	Used for decisions, not transactions

KEY IDEA · DECISIONS A PM MAKES HERE

- Which data must be exact (money, orders) and which can be seconds or minutes old (menus, ratings, rider dots).
- Which order states allow a cancel, and who pays.
- Dispatch rules: speed for the customer against fair earnings for riders.
- What the user sees while background work is still running.
- Which background job gets priority when the queue backs up.
- How often to update live data, trading smoothness against cost and battery.

IN THE INTERVIEW

“How would you design order tracking?” Don't draw twenty boxes. Say: “Orders and payments go in a SQL database because they must be exact. Rider locations, 75,000 a second at peak, go through a queue into an in-memory store and are pushed to the customer. Menus are cached. On a big night I'd protect rider dispatch first.”

Connect the dots

Four ideas ran through every case, product and technology topic here. A PM in this industry meets all four in the same week, often on the same rainy evening.

Thread	In the cases	In the products	In the technology
Density decides the economics	Gates in a few societies slowed a whole area's riders (Little's law); vendors short-supplying a store empty its shelves	Blinkit's mature stores; Swiggy's shared fleet; BigBasket's 20-home van route	Queue worker maths; cache hit rates that keep cost per order low at peak
The promise is the product	Cancellations rose when the arrival time jumped after payment; breaches traced to one step	Arrival time as a range; Bolt's speed from distance, not rushing	Live tracking at 75,000 pings a second; "show what is happening"
Money per order, then per customer	Swiggy's IPO story: contribution and payback; reservations monetised by restaurants, not diners	Platform fees, ads, owning stock; Instamart's breakeven; Zepto's young stores	Payments in a SQL database because money must be exact
More than one side to keep happy	Big brands and bakeries as vendors; restaurants' fear of no-shows	magicpin and ONDC as the restaurants' way out; the CCI report	Dark patterns: fees hidden from customers win a quarter, lose the regulator

BACK TO MUMBAI · A SKETCH ANSWER FOR THE MONSOON EVENING

Scope first. Goal: deliver as many orders as possible on time tonight without putting riders at risk. Who: customers, riders and restaurants in the flooded zones. Constraint: rider safety is not up for trade.

Next 30 minutes. (1) Make the promise honest before payment: wider arrival ranges calculated from live rider supply. (2) A rain fee shown up front on every order, funding a per-order rain incentive for riders who choose to stay online. (3) Hide restaurants more than about 4 km away and pause zones with waterlogged roads. (4) Batch orders from the same restaurant cluster so each ride carries two. (5) On the tech side, give the rider-dispatch queue priority over invoices and analytics.

What must not break. No incentive that pays riders to ride through water; no fee added after payment; no promise the fleet can't keep.

Goal (North Star): orders delivered on time per hour in the affected zones. **Safety checks (guardrails):** customer cancellations, rider incident reports, complaints about fees. **Trade-off:** fewer orders taken tonight, in exchange for promises kept and riders safe.

KEY IDEA · THE DELIVERY PM'S THREE QUESTIONS

1. Where do the minutes go, and which step can we change?
2. What does this order earn after the rider, the discount and the refund?
3. Which side (customer, restaurant, rider, brand) pays for this change?

IN THE INTERVIEW

In any delivery or quick-commerce case, break the time or the money into steps within the first three minutes: "A 12-minute order is assignment, picking, the ride and the gate", or "a ₹450 order earns about ₹145 and pays ₹85 to the rider". It shows you know how the machine works before you try to fix it.